

Peter Willendrup DTU Physics & ESS DMSC

McStas + McXtrace: Monte Carlo Ray-Tracers + GPU Using ISO-C Code-Generation and OpenACC #pragmas

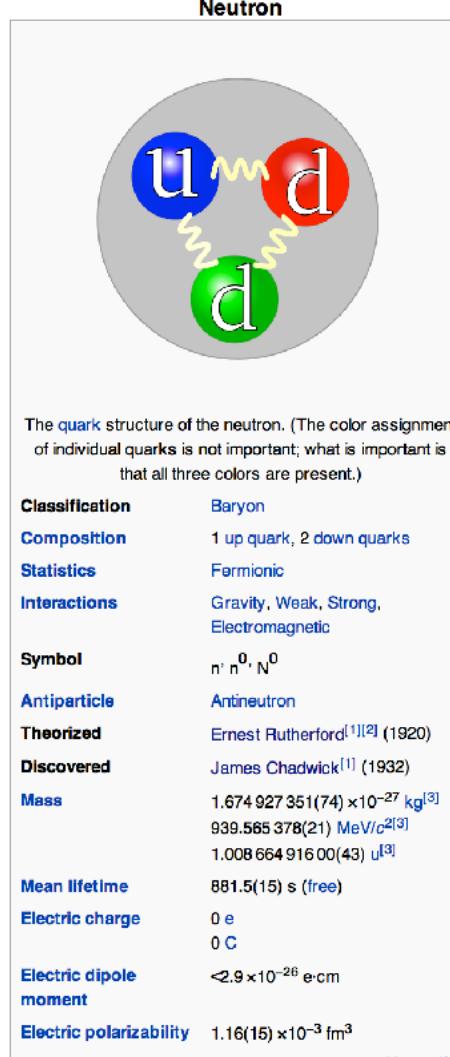


Agenda

- Background, science:  Keywords: Monte Carlo ray-traces for particle transport
 - Neutrons, X-rays and scattering techniques
 - What are the McStas and McXtrace codes used for?
- Background, code:
 - Technical foundations of the codes and why we chose OpenACC
 - Our timeline toward the GPU
 - What did we change for OpenACC and what does not (yet) work?
- GPU status:
 - How well (fast) does it work?
 - Simulation flow
 - What to investigate / improve?



The neutron and its basic properties



Life time: $\tau_{1/2} = 890s$

Mass: $m = 1.675 \times 10^{-27} \text{ kg}$

Charge: $Q = 0$

Spin: $s = \hbar/2$

Magnetic moment: $\mu/\mu_n = -1.913$

$$E = \frac{1}{2}mv^2 = \frac{\hbar^2k^2}{2m} \quad \lambda = 2\pi/k$$

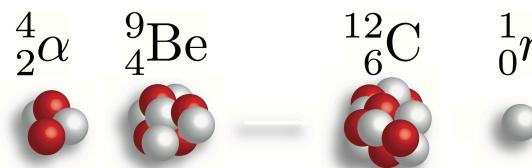
$$E = 81.81 \cdot \lambda^{-2} = 2.07 \cdot k^2 = 5.23 \cdot v^2$$

	Energy	Wavelength	n-Wavevector	Velocity	Frequency
cold neutrons:	$E = 1 \text{ meV}$ $E = 5 \text{ meV}$	$\lambda = 9.0446 \text{ \AA}$ $\lambda = 4.0449 \text{ \AA}$	$k = 0.6947 \text{ 1/\AA}$ $k = 1.5534 \text{ 1/\AA}$	$v = 437 \text{ m/s}$ $v = 978 \text{ m/s}$	$v = 0.2418 \text{ THz}$ $v = 1.2090 \text{ THz}$
thermal neutrons:	$E = 25 \text{ meV}$ $E = 50 \text{ meV}$	$\lambda = 1.8089 \text{ \AA}$ $\lambda = 1.2791 \text{ \AA}$	$k = 3.4734 \text{ 1/\AA}$ $k = 4.9122 \text{ 1/\AA}$	$v = 2187 \text{ m/s}$ $v = 3093 \text{ m/s}$	$v = 6.045 \text{ THz}$ $v = 12.090 \text{ THz}$

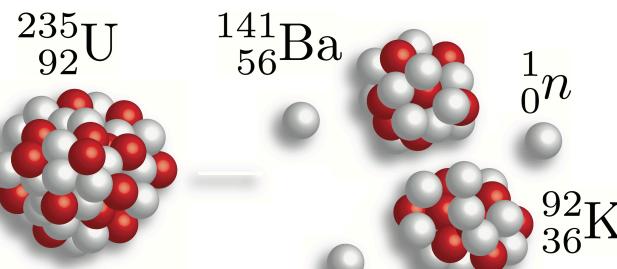
Subatomic particle discovered by Sir James Chadwick in 1932



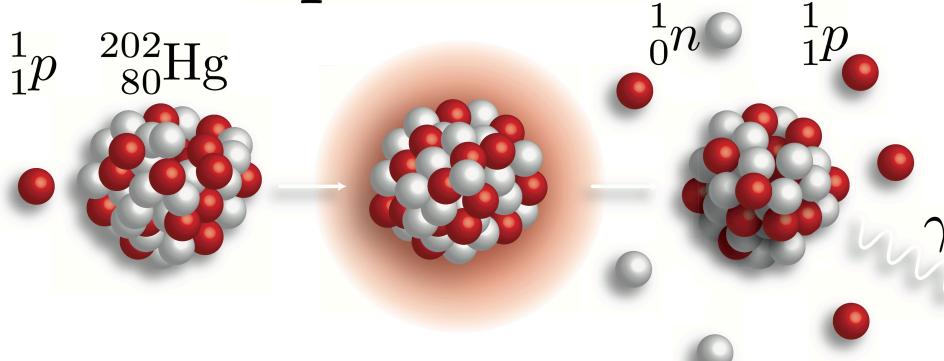
Chadwick



Fission

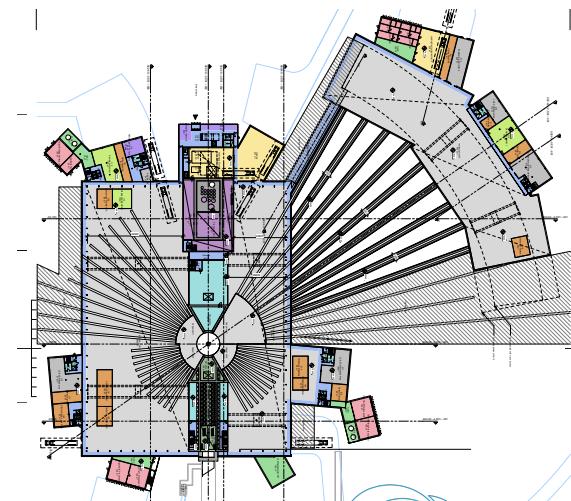
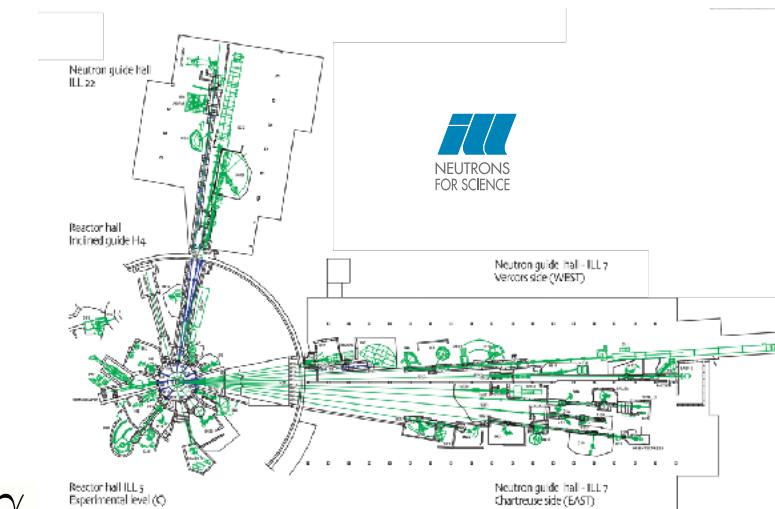
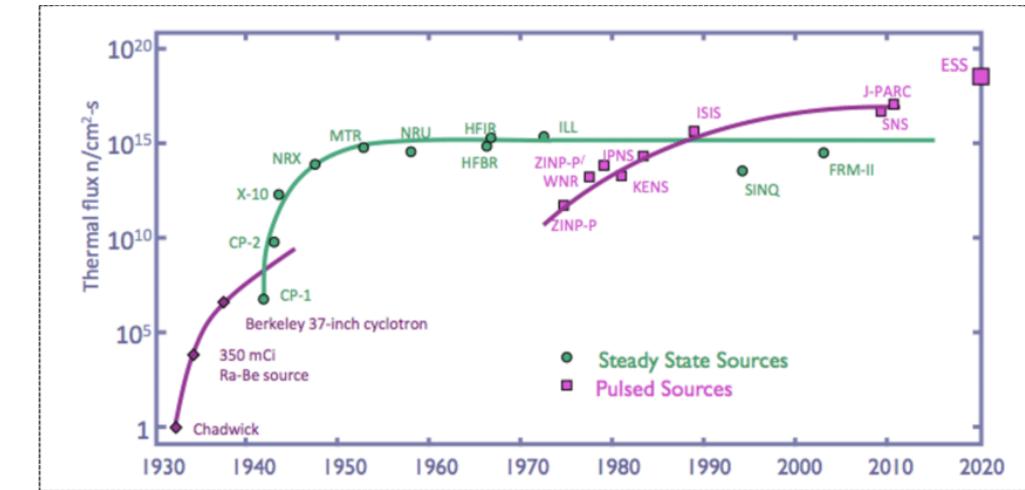


Spallation



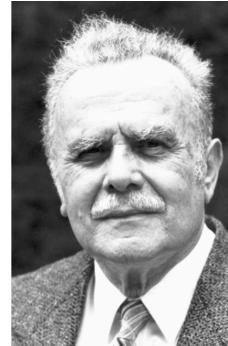
McXtrace

McStas



ess
EUROPEAN SPALLATION SOURCE

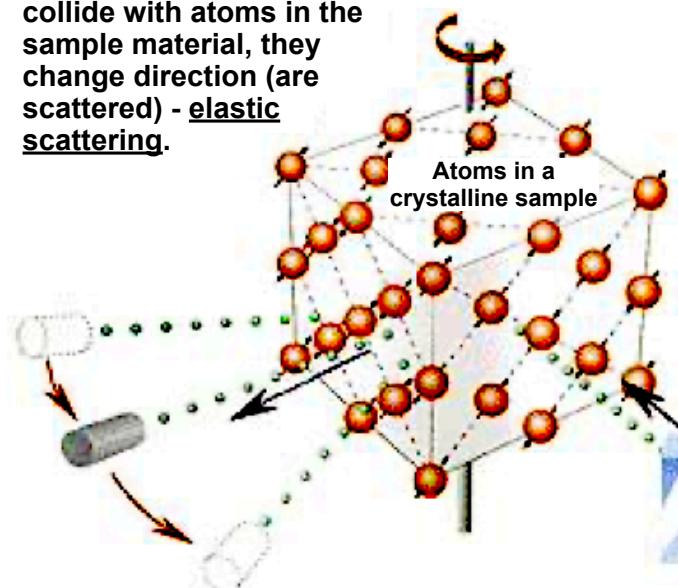
1994 Nobel Prize in Physics



'... In simple terms, Clifford G. Shull has helped answer the question of where atoms "are" and Bertram N. Brockhouse the question of what atoms "do".'

Diffraction

When the neutrons collide with atoms in the sample material, they change direction (are scattered) - elastic scattering.

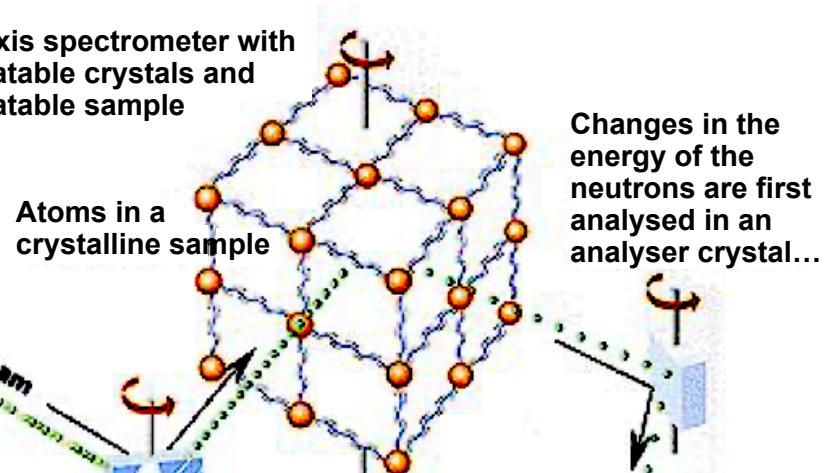


Detectors record the directions of the neutron and a diffraction pattern is obtained.

The pattern shows the positions of the atoms relative to one another.

Spectroscopy

3-axis spectrometer with rotatable crystals and rotatable sample



When the neutrons penetrate the sample they start or cancel oscillations in the atoms. If the neutrons create phonons or magnon they themselves lose the energy these absorb - inelastic scattering.

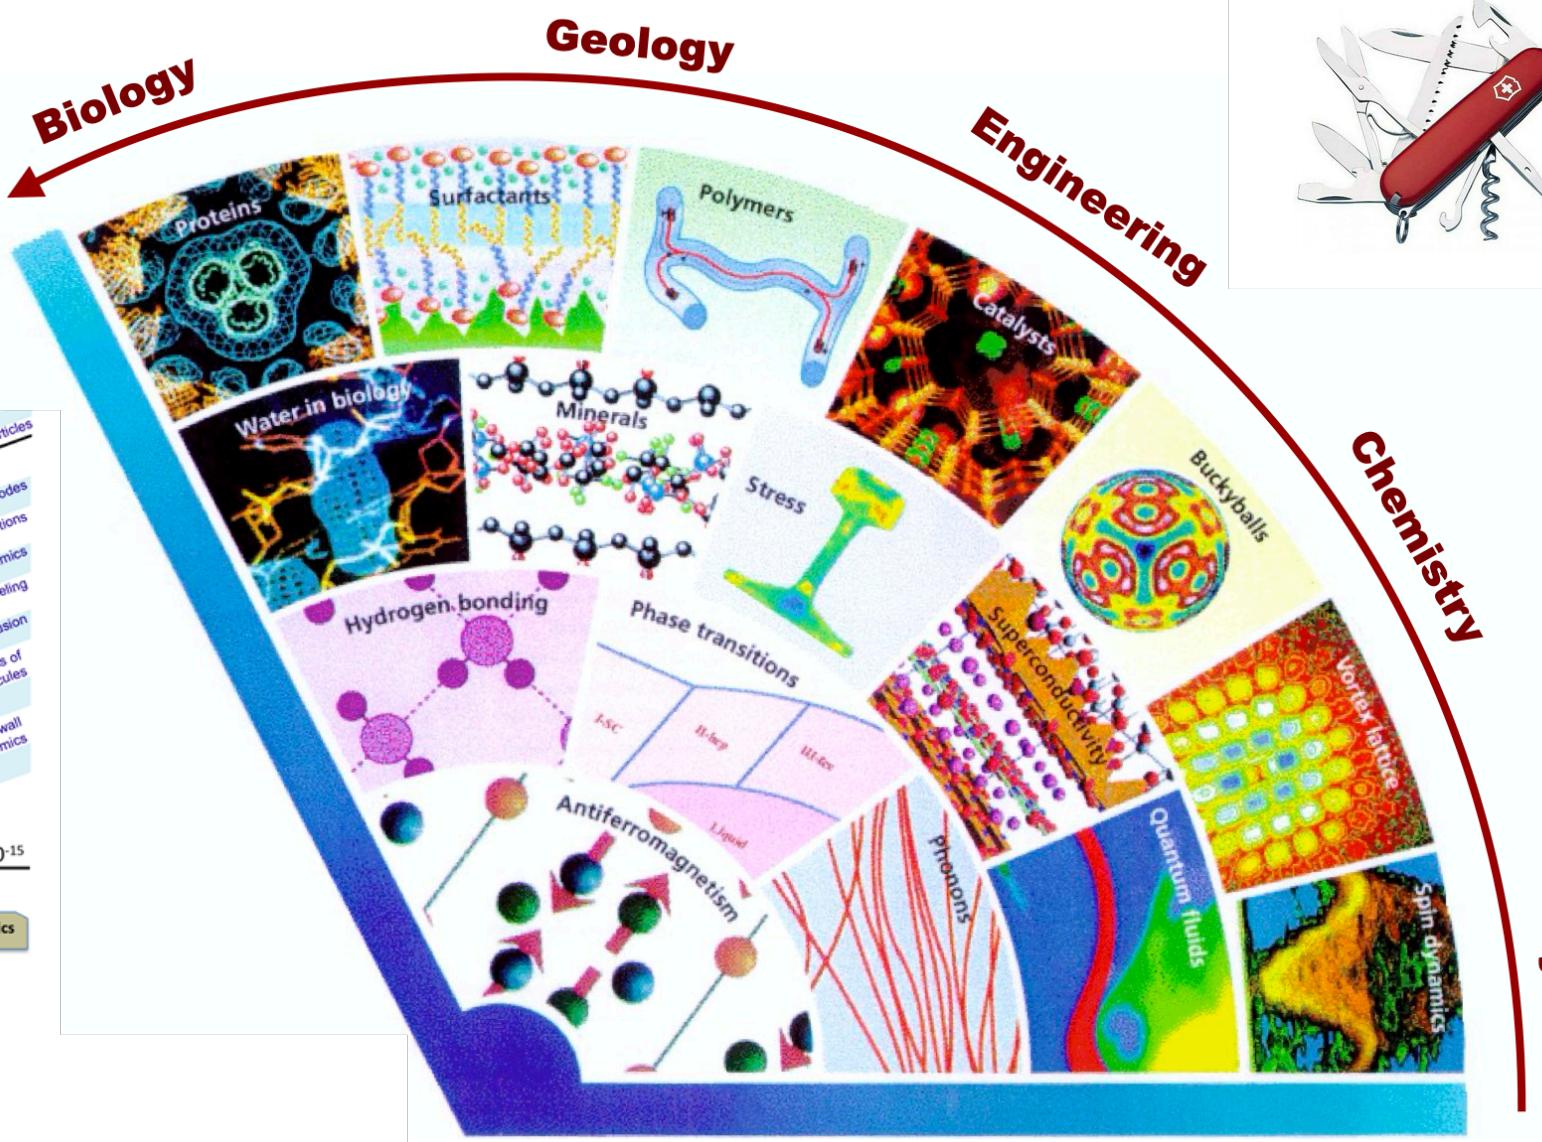
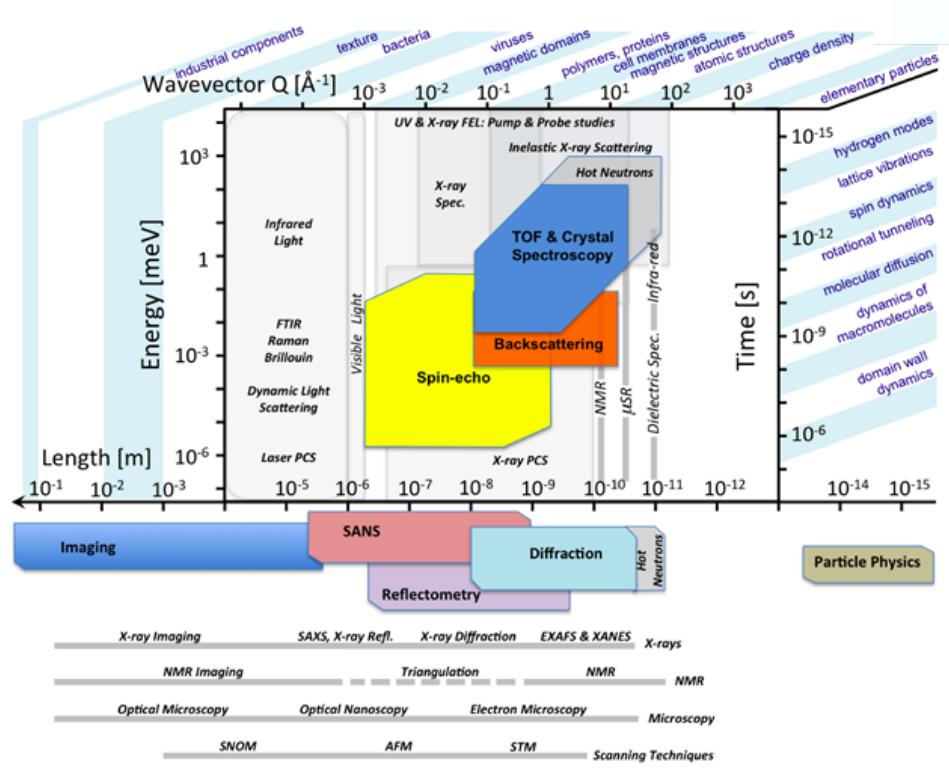
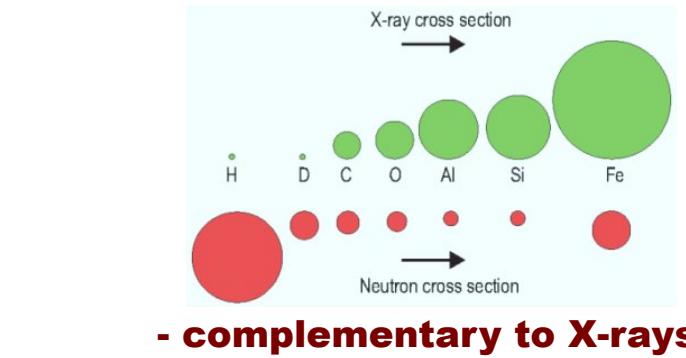
Changes in the energy of the neutrons are first analysed in an analyser crystal...

.. and the neutrons are counted in a detector.

Photo from the Nobel Foundation archive.
Bertram N. Brockhouse
Prize share: 1/2

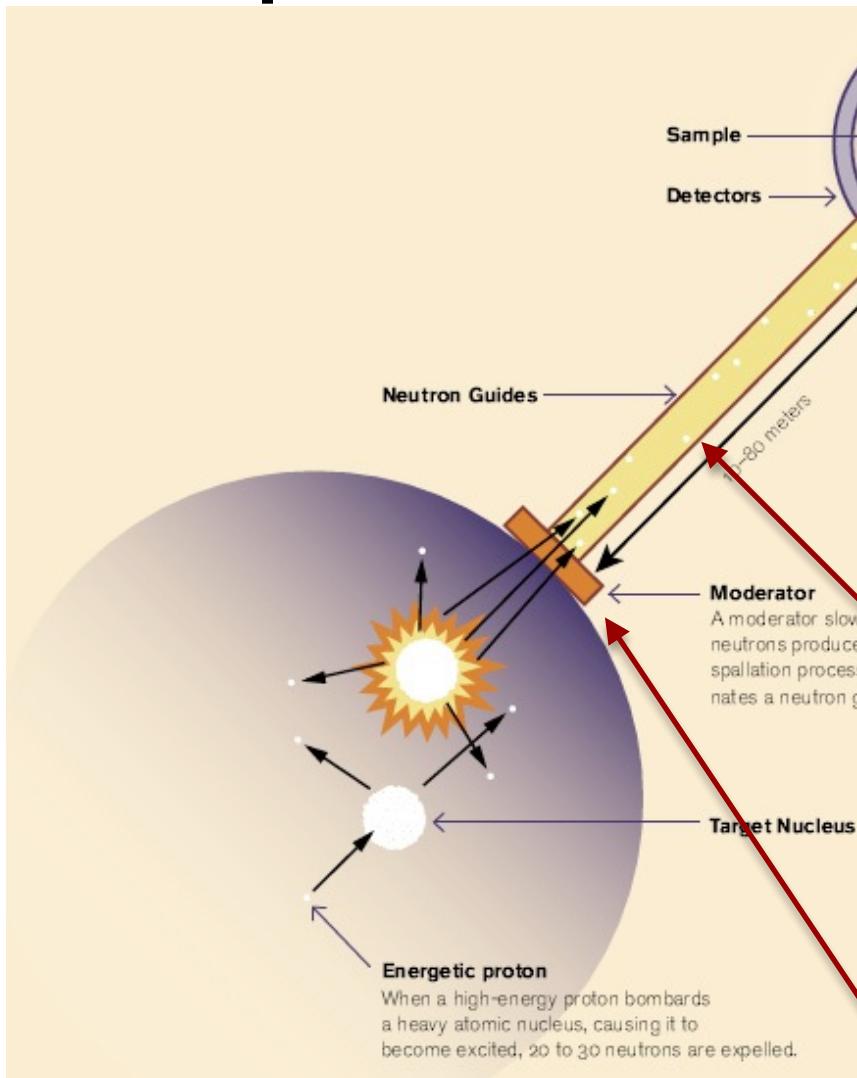
Photo from the Nobel Foundation archive.
Clifford G. Shull
Prize share: 1/2

The neutron is a multidisciplinary probe for structure and dynamics of condensed matter systems - ‘Swiss army knife’



Components of neutron instruments

McStas



The diagram illustrates the basic components of a neutron instrument. It starts with an **Energetic proton** hitting a **Target Nucleus**, which produces a group of neutrons. These neutrons pass through a **Moderator** (represented by a purple sphere) to slow them down. The moderated neutrons then travel along **Neutron Guides** (yellow tubes) for approximately 1-80 meters before reaching a **Sample** (represented by a purple cylinder). After interacting with the sample, the neutrons are detected by **Detectors**.

Detectors are “monitors” in McStas. Mostly they act as “perfect probes” and can be positioned thought your instrument gathering 1D/2D/ event lists...

The sample:
Crystalline, powders, liquids, micelles, structures to image, inelastic features like phonons...

Neutron optics include things like:

- Mirrors and guides
- Collimators and slits
- Diskchoppers, Fermi ch and velocity selectors
- Monochromators/Analysers

In McStas the moderator is the “source”

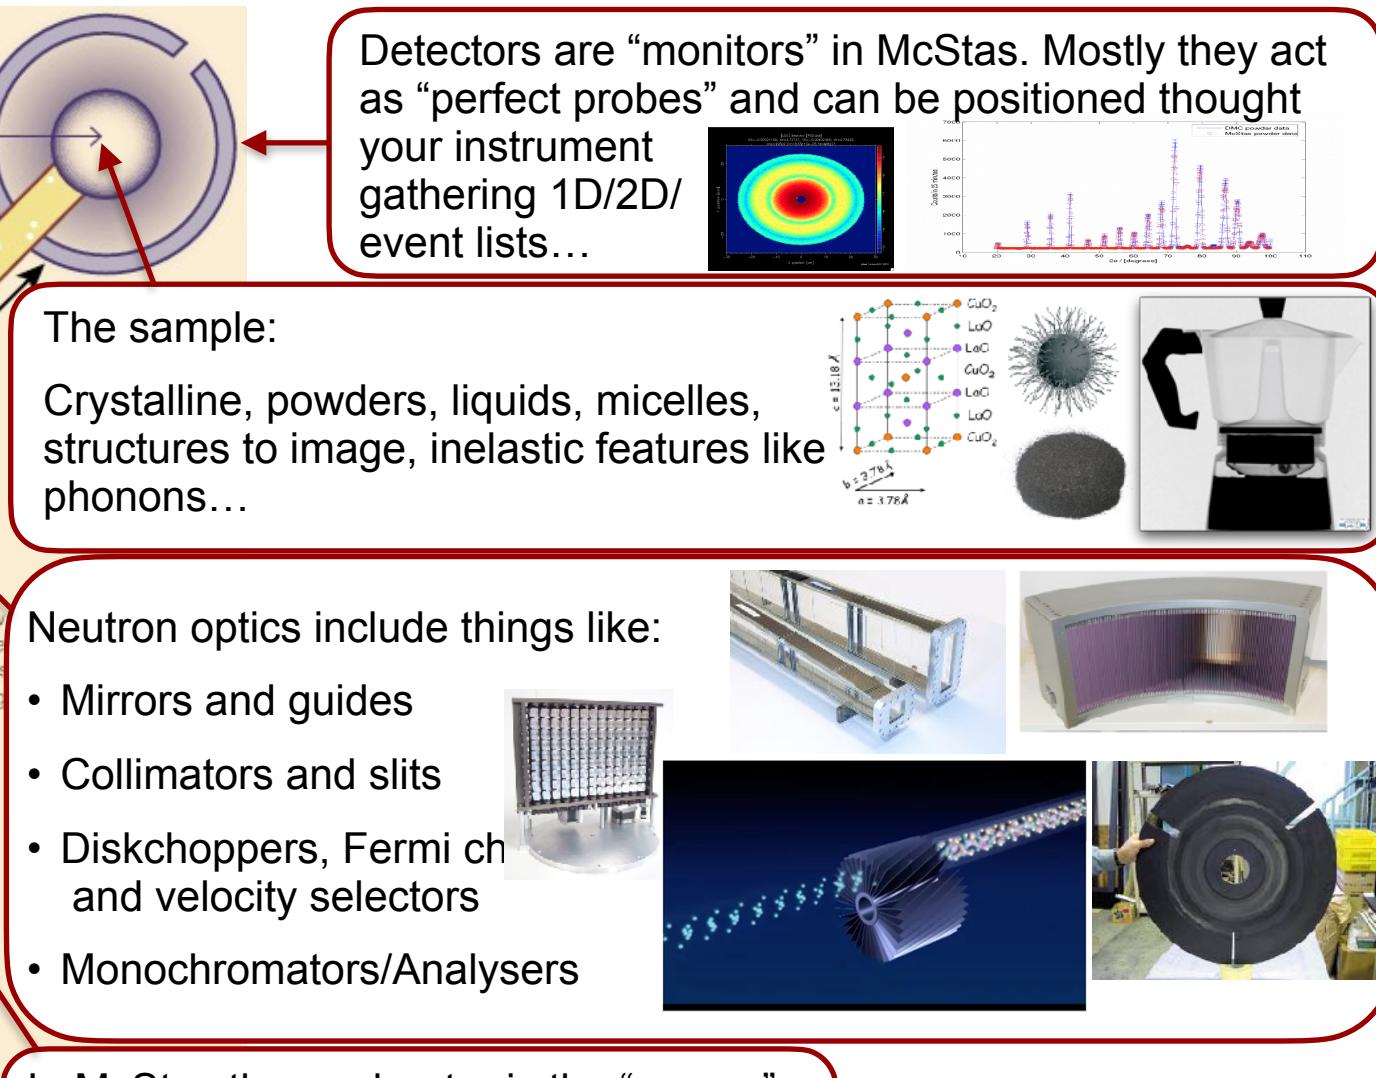


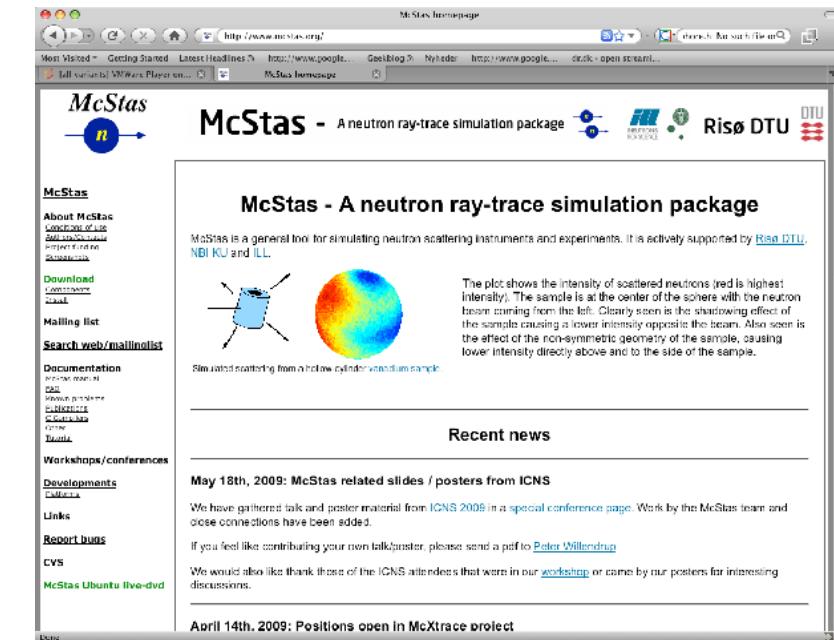
Figure showing neutron scattering data and neutron optics components. The top right shows a 2D intensity map of a sample, a 1D intensity profile plot, and a 2D plot of OMA powder data. Below these are crystallographic unit cell diagrams for CuO₂ and LaO, and a photograph of a neutron monochromator. The bottom right shows various neutron optics components: a curved mirror, a rectangular collimator, a disk chopper, and a large black monochromator wheel.

McStas Introduction

- Flexible, general simulation utility for neutron scattering experiments.
- Original design for **Monte carlo Simulation of triple axis spectrometers**
- Developed at DTU Physics, ILL, PSI, Uni CPH, ESS DMSC
- V. 1.0 by K Nielsen & K Lefmann (1998) RISØ
- Currently ~6 people on joint McStas-McXtrace team but only 2 full time, based at DTU



GNU GPL license
Open Source



The screenshot shows the McStas homepage. On the left is a sidebar with links like "About McStas", "Documentation", "Workshops/conferences", and "Recent news". The main content area features a heading "McStas - A neutron ray-trace simulation package" and a sub-section "McStas - A neutron ray-trace simulation package". It includes a plot titled "Simulated scattering from a hollow cylinder vanadium sample" showing intensity distribution. Below the plot, there's a section for "Recent news" with entries for May 18th, 2009, and April 14th, 2009.

Project website at

<http://www.mcstas.org>

mcstas-users@mcstas.org mailinglist

McXtrace - since jan 2009 similar for X-rays

Main Page – McXtraceWiki

Most Visited ▾ Getting Started Latest Headlines ▾ http://www.google... Geekblog ▾ Nyheder http://www.google... dr.dk ▾ open streami... Log in / create account

[article](#) [discussion](#) [edit](#) [history](#)

McXtrace

Main Page [edit]

McXtrace

McXtrace - Monte Carlo Xray ray-tracing is a joint venture by

Risø DTU DTU ESRF JJ X-RAY

Funding from NABIIT, DSF and the above parties.

Our code will be based on technology from 

For information on our progress, please subscribe to our user mailinglist.
<mailto:webmaster@mcxtrace.org>

This page was last modified 13:15, 25 February 2009. This page has been accessed 2,049 times. Privacy policy About McXtraceWiki Disclaimers Powered By MediaWiki

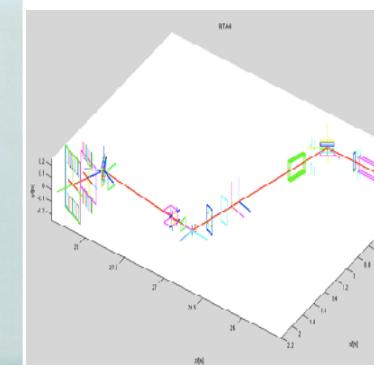
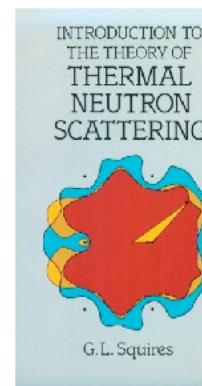
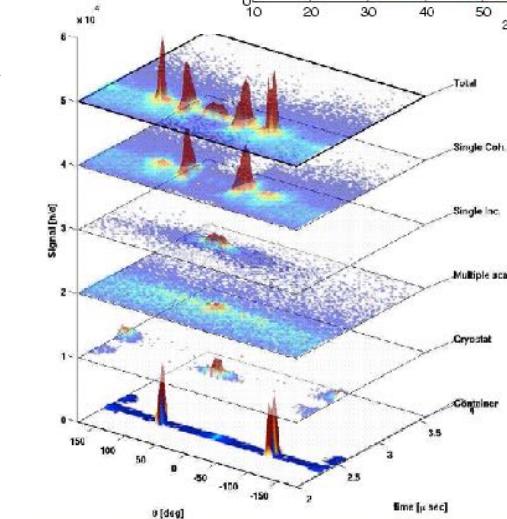
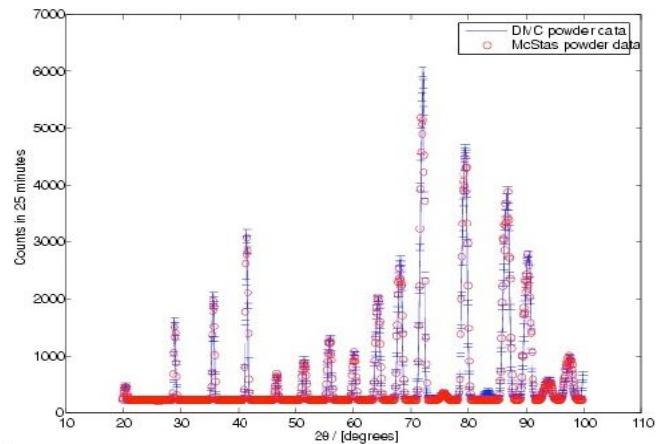
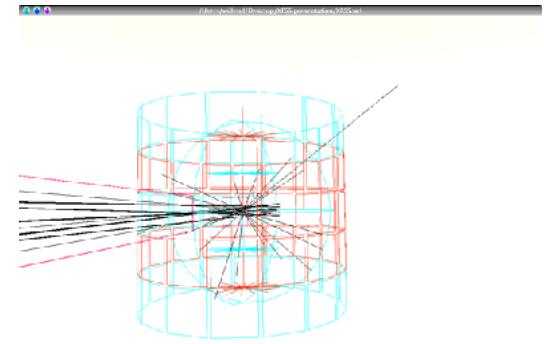
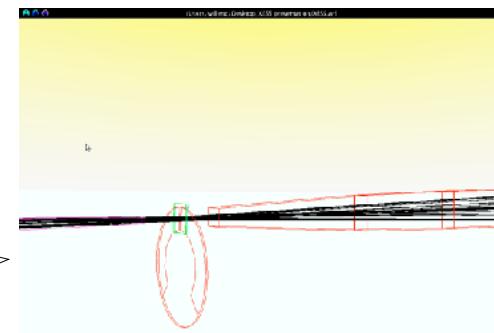
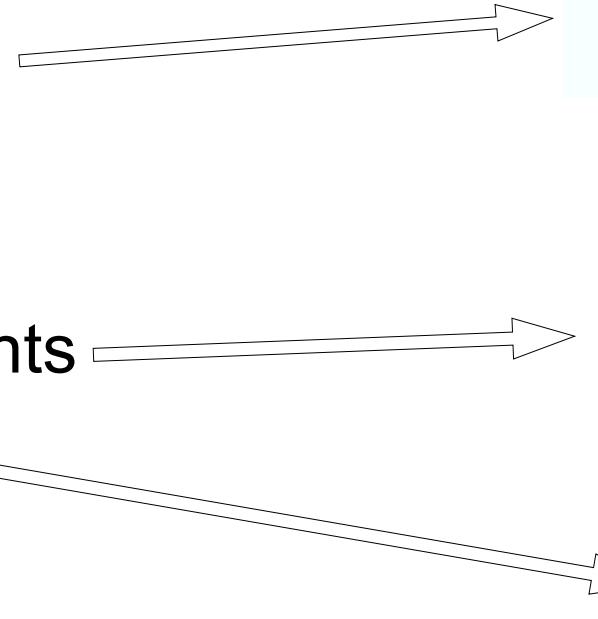
- Synergy, knowledge transfer, shared infrastructure and codebase on GitHub



What is McStas used for?

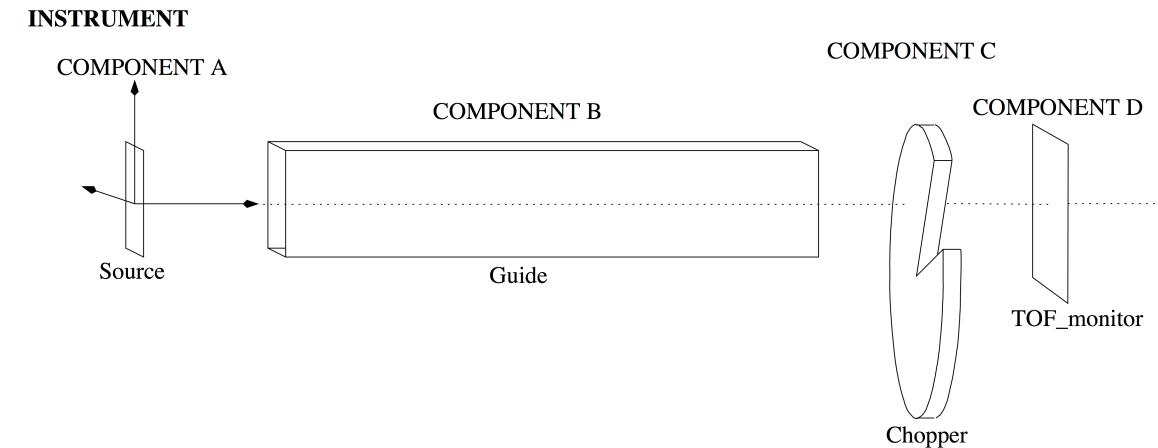
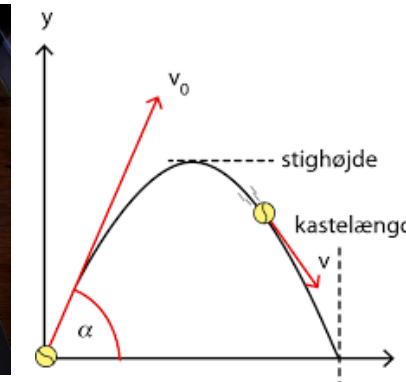
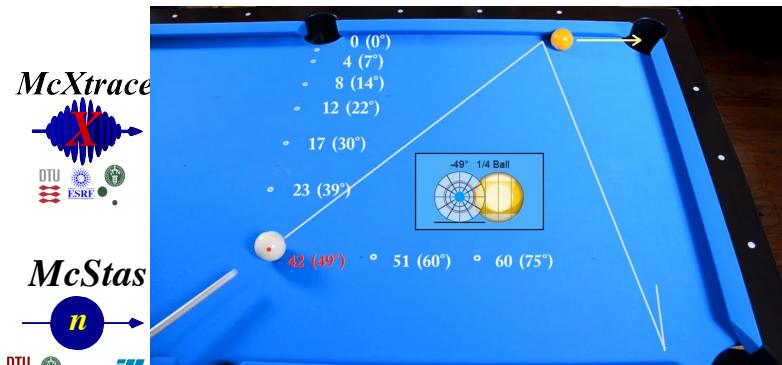
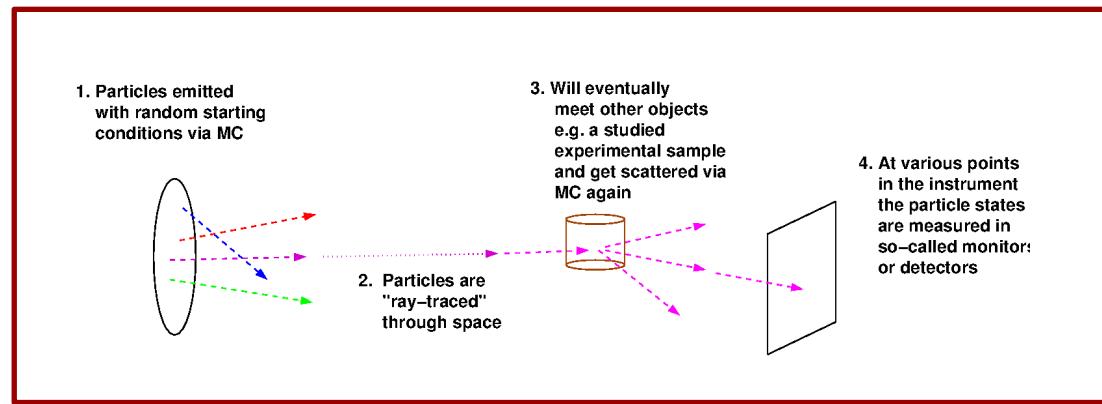
- Instrumentation
- Planning
- Construction
- Virtual experiments
- Data analysis
- Teaching

(KU, DTU)



McStas and McXtrace

Monte Carlo ray-tracers



The "tool layer" consists of programs manipulated by the McStas user:

mcgui, graphical user interface

mcplot, visualize histogram outp.

mcdisplay, visualize instrument

mcgui is used to assemble an instrument file, which is taken over by the McStas system

DEFINE INSTRUMENT Example(Param1=1, string Param2="two", ...)

COMPONENT A = Source(Parameters...)
AT (0, 0, 0) ABSOLUTE

COMPONENT B = Guide(Parameters...)
AT (0, 0, 1) RELATIVE A

COMPONENT C = DiskChopper(Parameters...)
AT (0, 0, 1) RELATIVE B

COMPONENT D = TOF_monitor(Parameters, filename="Tof.dat")
AT (0, 0, Param1) RELATIVE PREVIOUS

DSL

"Instrument file"

Source.comp – c-code

Guide.comp – c-code

DiskChopper.comp – c-code

TOF_monitor.comp – c-code

Component library

Random numbers

I/O

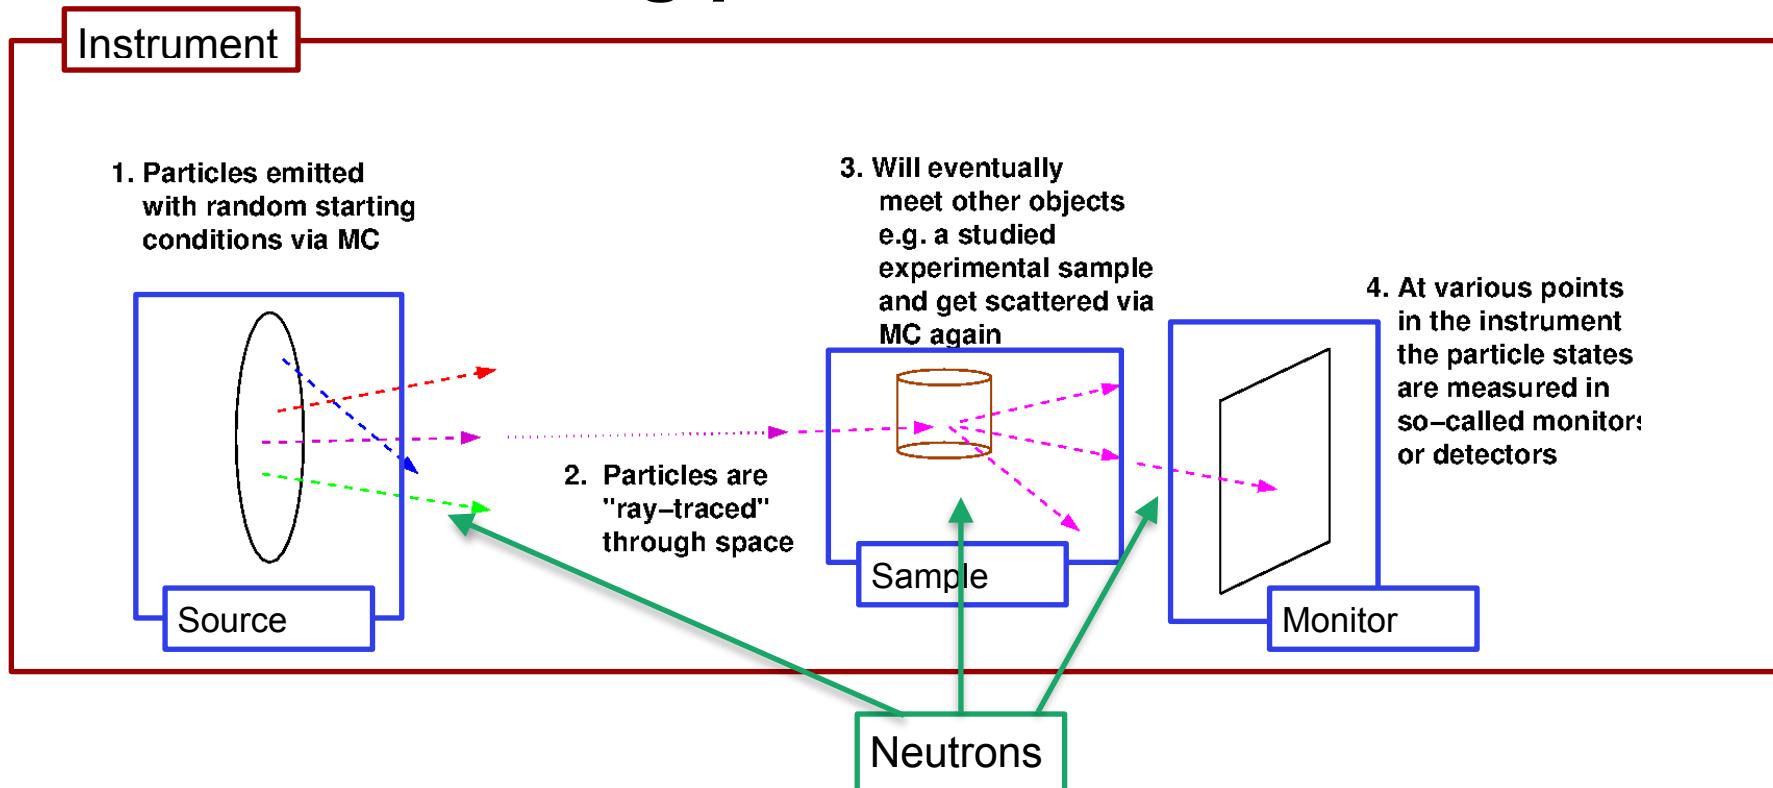
Physical const.

"Kernel and runtime c-code"

The McStas system generates an "ISO C file" and an executable from instrument file and c-codes

The simulation executable produces data output which can be visualized using the mcplot and mcdisplay tools

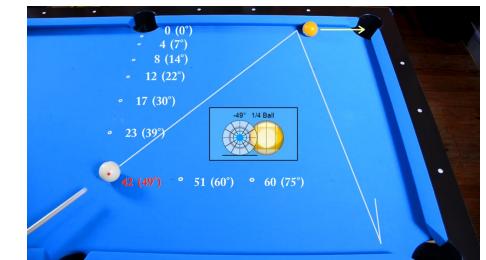
In the big picture, McStas is this...



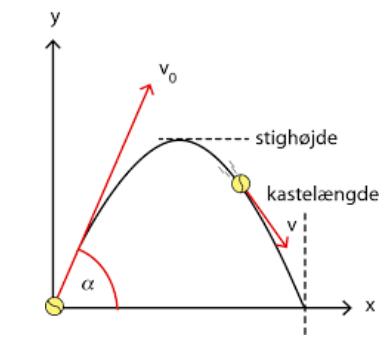
The instrument defines our “lab coordinate system”

The components define devices or features available in our instrument - they have different function

Neutron particles are passed on from one component to the next, changing state under way

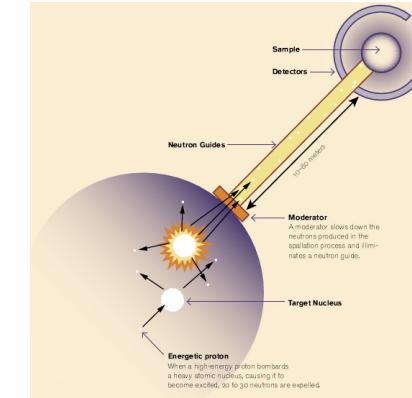


Classical Newtonian mechanics
(independent, particles though...)

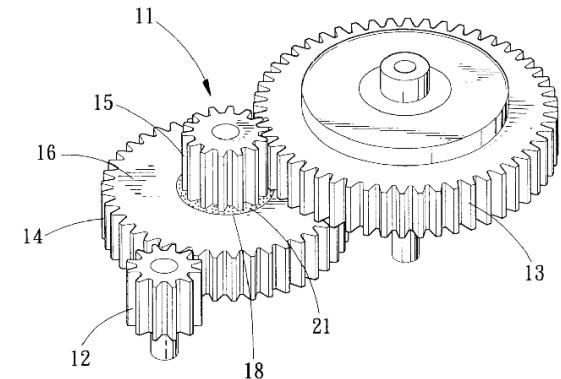


McStas tech overview

- Portable code (Unix/Linux/Mac/Windoze)
 - On the CPU-side, ran on everything from iPhone to 1000+ node cluster, intel, Alpha, PA-RISC etc.
- 'Component' files (~100) inserted from library
 - Sources
 - Optics
 - Samples
 - Monitors
 - If needed, write your own comps - **many are USER developments ~200-line “physicist” codes**



- DSL + ISO-C code-gen. (compiler technology / LeX+Yacc)
 - Simple Instrument language  ISO C
- Component codes realizing beamline parts (including user contribs)
- Library of common functions
 - I/O
 - Random numbers
 - Physical constants
 - Propagation
 - Precession in fields
 - ...



User experience:

- Write instrument
- Launch simulation (generates binary and runs simulation)
- Look at output data

<https://www.openhub.net/p/mccode>

Stats and information on the codebase

Total Lines :	2,132,114	Code Lines :	1,832,884	Percent Code Lines :	86.0%
Number of Languages :	23	Total Comment Lines :	203,852	Percent Comment Lines :	9.6%
		Total Blank Lines :	95,378	Percent Blank Lines :	4.5%

Language Breakdown

Language	Code Lines	Comment Lines	Comment Ratio	Blank Lines	Total Lines	Total Percentage
Postscript	901,336	9,202	1.0%	948	911,486	<div style="width: 42.8%;"><div style="width: 42.8%;"></div></div> 42.8%
AMPL	356,420	27,344	7.1%	4,012	387,776	<div style="width: 18.2%;"><div style="width: 18.2%;"></div></div> 18.2%
Fortran (Free-format)	216,230	86,604	28.6%	34,678	337,512	<div style="width: 15.8%;"><div style="width: 15.8%;"></div></div> 15.8%
C	126,112	31,224	19.8%	20,646	177,982	<div style="width: 8.3%;"><div style="width: 8.3%;"></div></div> 8.3%
JavaScript	100,016	1,116	1.1%	2,596	103,728	<div style="width: 4.9%;"><div style="width: 4.9%;"></div></div> 4.9%
Python	45,988	33,214	41.9%	18,278	97,480	<div style="width: 4.6%;"><div style="width: 4.6%;"></div></div> 4.6%
TeX/LaTeX	37,440	4,208	10.1%	5,940	47,588	<div style="width: 2.2%;"><div style="width: 2.2%;"></div></div> 2.2%
shell script	11,870	2,836	19.3%	2,954	17,660	<div style="width: 0.8%;"><div style="width: 0.8%;"></div></div> 0.8%
CMake	10,746	2,522	19.0%	2,492	15,760	<div style="width: 0.7%;"><div style="width: 0.7%;"></div></div> 0.7%
HTML	6,712	0	0.0%	238	6,950	<div style="width: 0.3%;"><div style="width: 0.3%;"></div></div> 0.3%
XML	4,486	58	1.3%	68	4,612	<div style="width: 0.2%;"><div style="width: 0.2%;"></div></div> 0.2%
Matlab	3,362	782	18.9%	550	4,694	<div style="width: 0.2%;"><div style="width: 0.2%;"></div></div> 0.2%
Perl	3,120	3,392	52.1%	48	6,560	<div style="width: 0.3%;"><div style="width: 0.3%;"></div></div> 0.3%
C++	2,968	448	13.1%	696	4,112	<div style="width: 0.2%;"><div style="width: 0.2%;"></div></div> 0.2%
TypeScript	2,718	12	0.4%	312	3,042	<div style="width: 0.1%;"><div style="width: 0.1%;"></div></div> 0.1%
NSIS	898	316	26.0%	380	1,594	<div style="width: 0.1%;"><div style="width: 0.1%;"></div></div> 0.1%
DOS batch script	696	236	25.3%	108	1,040	<div style="width: 0.0%;"><div style="width: 0.0%;"></div></div> 0.0%
CSS	662	8	1.2%	96	766	<div style="width: 0.0%;"><div style="width: 0.0%;"></div></div> 0.0%
Make	344	136	28.3%	122	602	<div style="width: 0.0%;"><div style="width: 0.0%;"></div></div> 0.0%
Fortran (Fixed-format)	340	0	0.0%	122	462	<div style="width: 0.0%;"><div style="width: 0.0%;"></div></div> 0.0%
IDL/PV-WAVE/GDL	232	134	36.6%	58	424	<div style="width: 0.0%;"><div style="width: 0.0%;"></div></div> 0.0%
Vim Script	160	36	18.4%	14	210	<div style="width: 0.0%;"><div style="width: 0.0%;"></div></div> 0.0%
R	28	24	46.2%	22	74	<div style="width: 0.0%;"><div style="width: 0.0%;"></div></div> 0.0%
Totals	1,832,884	203,852		95,378	2,132,114	

Project Summary : Factoids

Mature, well-established codebase

Large, active development team

Stable Y-O-Y development activity

Very few source code comments

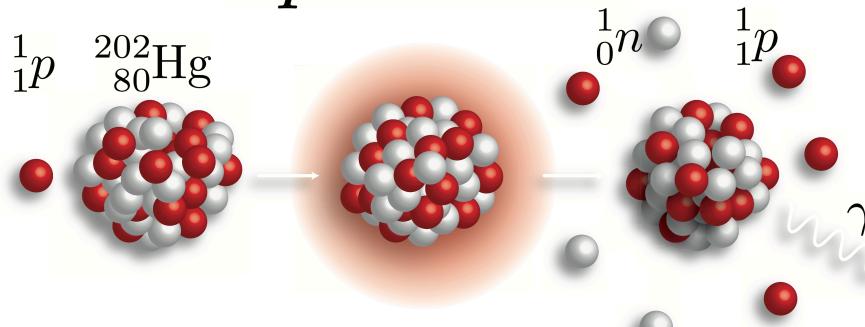
McCode is written mostly in AMPL.



Nope, that's our DSL and grammar. :-) Which is close to "English".

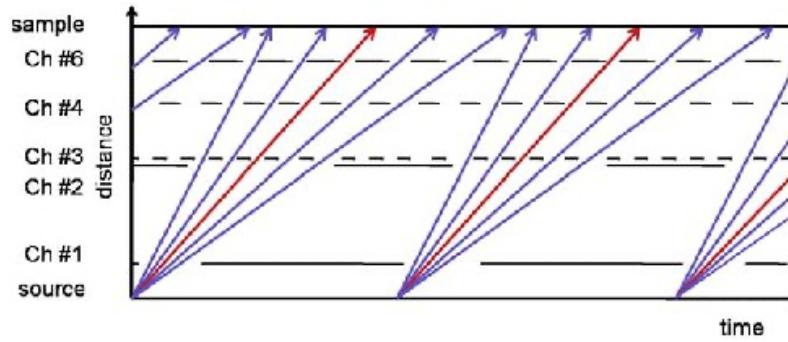
If it is so bloody good and “final”, why bother looking at GPU?

Spallation



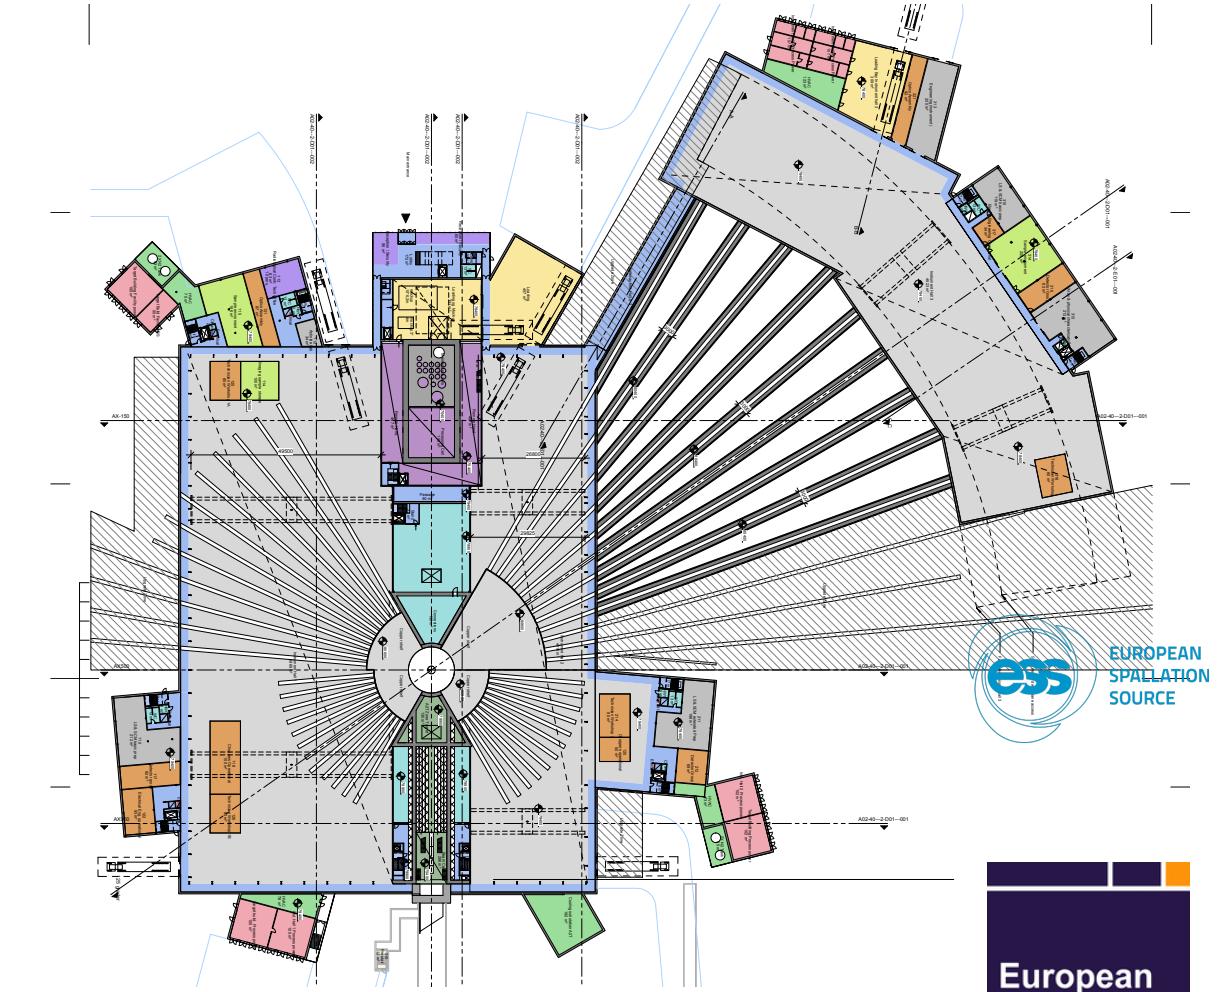
Next-generation, long-pulse spallation facilities are complex to construct and model since they use

- ‘rep-rate multiplication’
- event-mode experiments



At reactor and short-pulse, the **McStas run-time** keeps up with experimental time, **not the the case at ESS...**

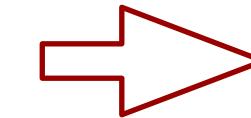
A ~2 order of magnitude would be great - and we believe we found it!



Even more so for e.g. X-ray Free-Electron Lasers...

GPU-computing in McStas - McXtrace

- LOTS of slowish processor-cores
- Limited bandwidth in CPU / GPU exchange



Perfect for highly-parallel calculations that have compact input / output, e.g. McStas.
(embarrassingly parallel)

- Main HW providers today are  **nVIDIA** ,  to some extent 

- Software-frameworks

- CUDA (NVIDIA-specific)

- Accessible in C, ++, Fortran, shared library

- OpenCL (hw-agnostic)

- C-like programming language with own syntax

- OpenACC (almost NVIDIA-specific, but extending

- #pragma pre-compiler mechanism, accessible in C, ++, Fortran

- OpenMP is picking up GPU-support...

- Intel oneAPI

- Claims to be a unified approach to CPU/GPU/MPI/... I didn't get around to really try it yet... ;-)


nVIDIA

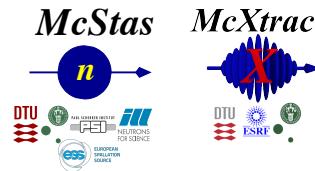
 to some extent **intel**

Prototype made by Brian Vinter students in 2010:
 * Good speed-up of ~2 orders found
 * Considered specific problem only
 * Was float-only solution in hw at the time
 * Would be too invasive on the code-base

Prototype made by Emmanuel Farhi + Post Doc in 2014:
 * Would be too invasive on the code-base
 * Effectively duplication of algorithm-code for CPU/GPU version
 * Lots of locks / mutex'es needed placement everywhere
 * Limited speedups found, e.g. ~ 15% if using GPU for RNG

What we eventually have gone for, with good results!

Main events on timeline of road toward GPU

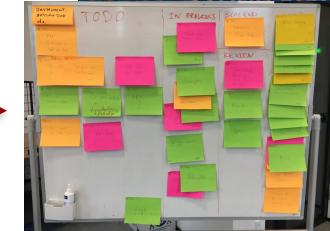


2017: E. Farhi
initial cogen
modernisation

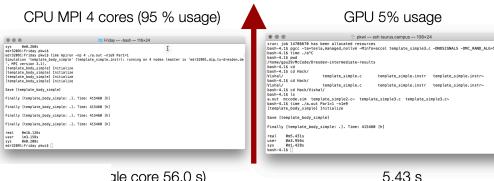
Fall 2018 onwards:
J. Garde further cogen
modernisation and
restructuring

October 2019 onwards:
J. Garde & P. Willendrup:
New RNG, test system, multiple
functional instruments.

January 2020:
One-week local
hackathon @ DTU
with McCode & RAMP teams



McStas / McXtrace instrument simulation



nVIDIA mentor: Vishal Metha

hackathon org.: **HZDR**
Guido Juckeland
HELMHOLTZ ZENTRUM DRESDEN ROSSENDORF

March 2018:
Participation at
Dresden Hackathon.
1st “null” instrument
prototype runs.

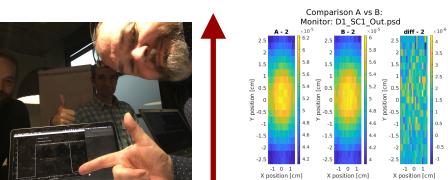


nVIDIA mentor: Christian Hundt

hackathon org.:
Sebastian Von Alfthan



October 2019:
Participation at Espoo
Hackathon. First meaningful
data extracted. Work on
cogen and realising we need
another RNG.



February 2020:
First release
McStas **3.0beta**
with GPU
support was
released
to the public

McStas heading for the GPU... numbers from November 2019

9 instruments fully ported, also realistic ones like PSI_DMC

(Aug 2020: 99 instrs)

10-core MPI run,
 $1e9$ in 200 secs

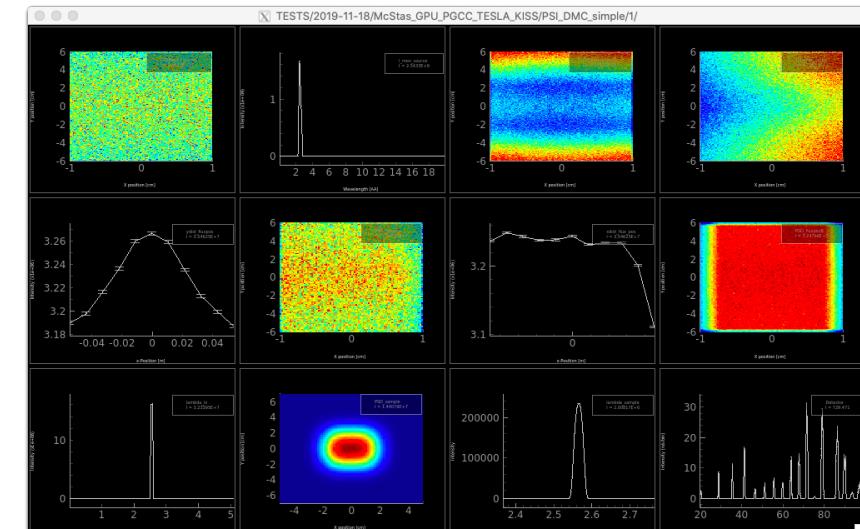
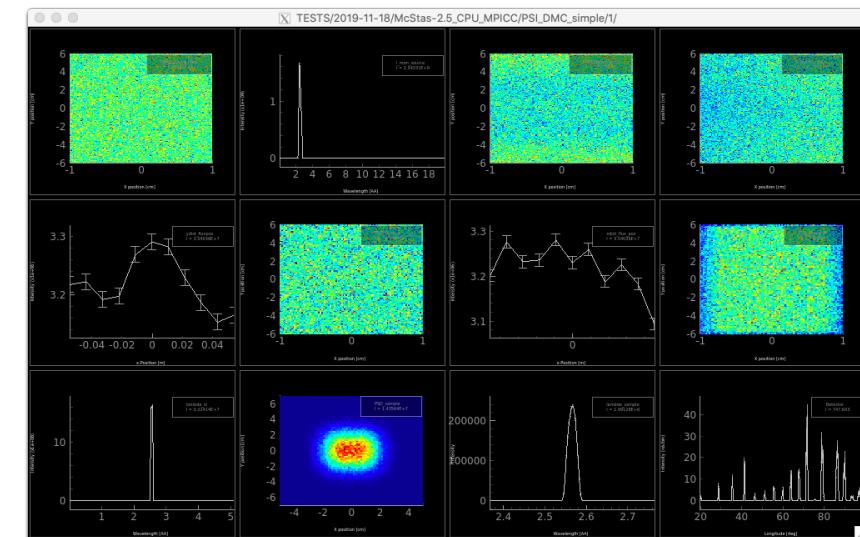


(1-core run,
 $1e9$ would be
2000 secs)

VS.

~ i.e. 2 orders of magnitude wrt. a single, modern CPU core

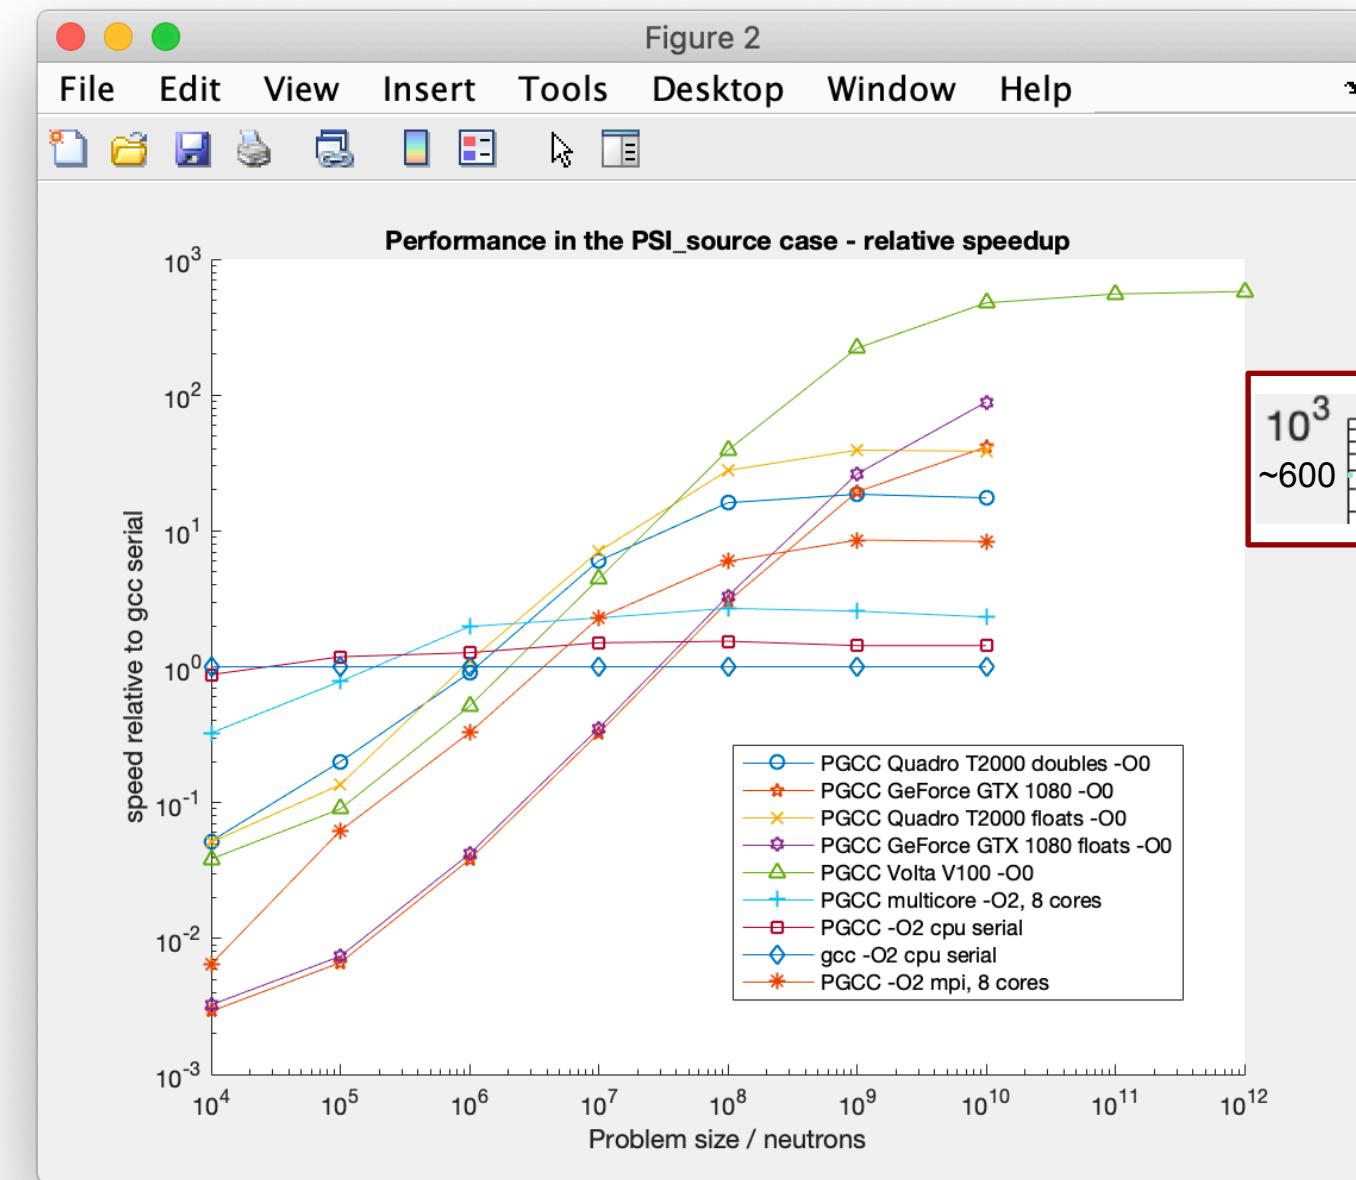
- If problem has the right size / complexity, GPU via OpenACC is great!



McStas heading for the GPU... November 2019 - first good look at performance.

Idealised instrument
with source and monitor
only - i.e. without any
use of the ABSORB
macro.

(Likely a good indication
of maximal speedup
achievable.)



Speedup

Looks like a factor of ~600



V100 execution speedups renormalised to wall-clock of single-core gcc standard simulation,

**V100 run is
600 times faster
than a single-
core i7 CPU run**

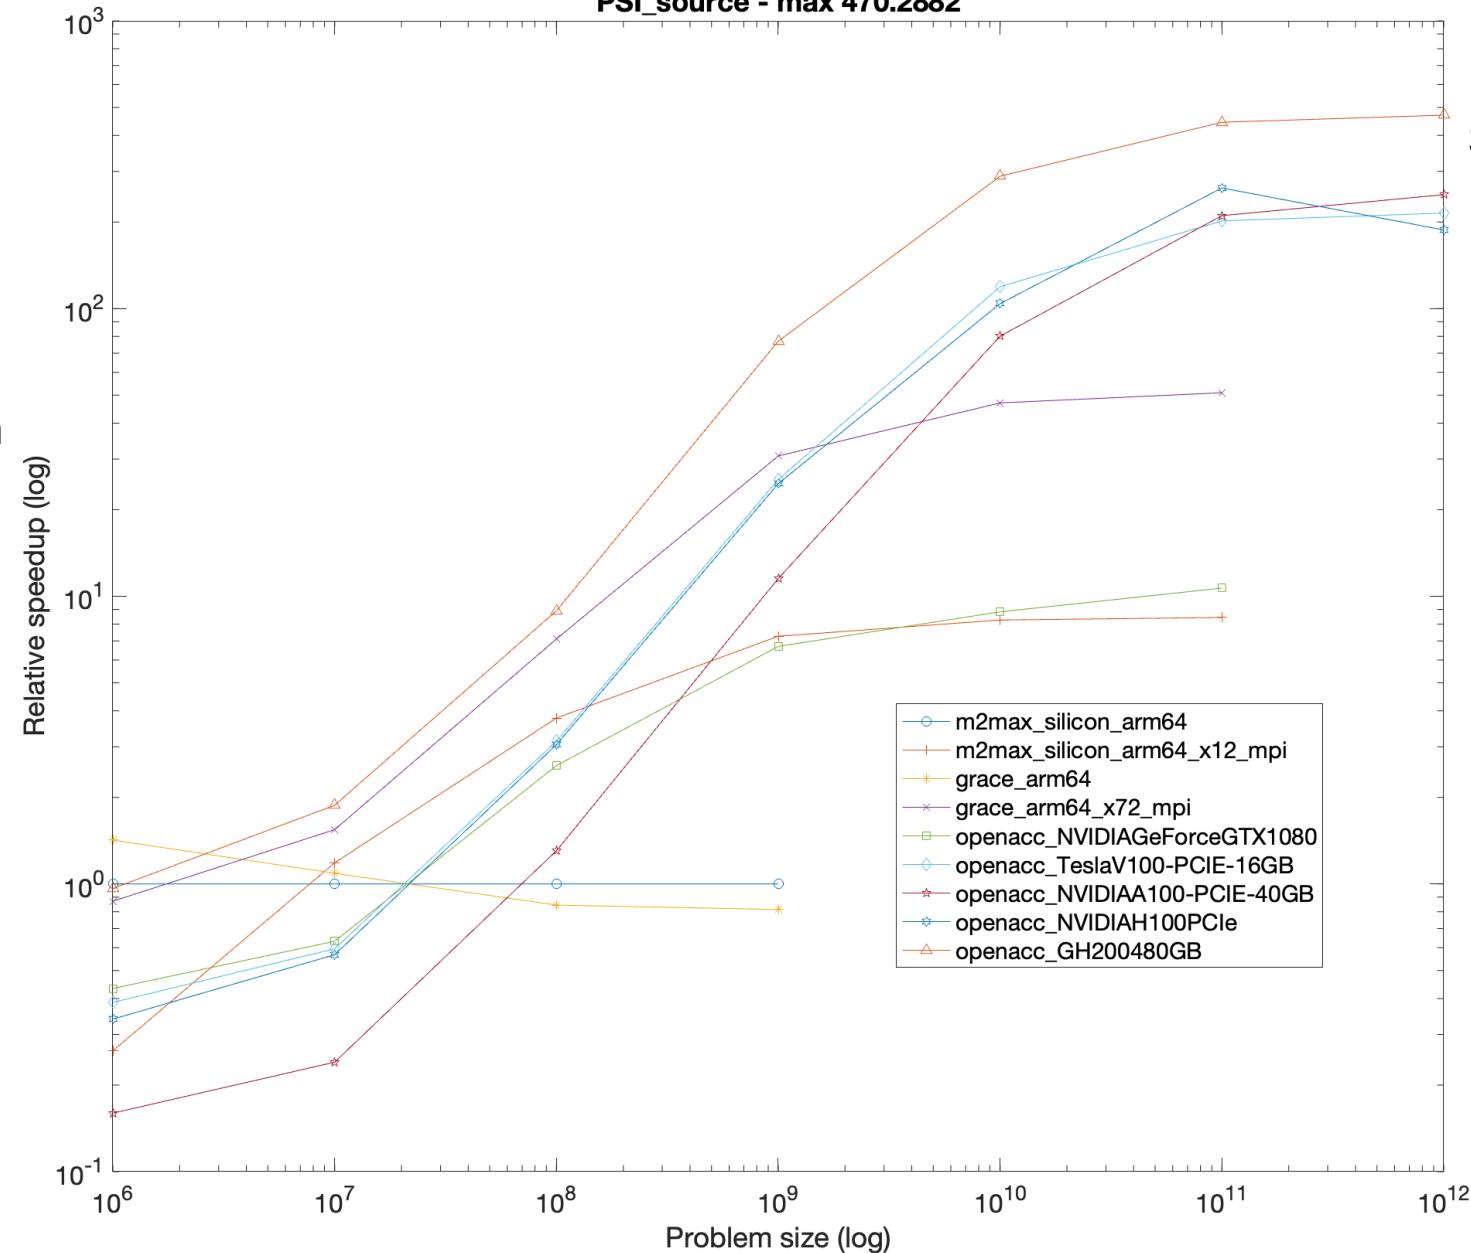
2025 measurement vs. arm64 cores in MacBook Pro / Grace Hopper

Idealised instrument
 with source and monitor
 only - i.e. without any
 use of the ABSORB
 macro.

(Likely a good indication
 of maximal speedup
 achievable.)

Speedup

Factor of ~500 GH200 vs
 single arm64 core



The neutron and “state-flags” in the instrument

v2.5: Global variables

```
double x, y, z, vx, vy, vz, t, sx, sy, sz, p;      double flag;
```

v3.0: particle struct, including any USERVARS like flag.

```
struct _struct_particle {
    double x,y,z; /* position [m] */
    double vx,vy,vz; /* velocity [m/s] */
    double sx,sy,sz; /* spin [0-1] */
    unsigned long randstate[7]; ←
    double t, p; /* time, event weight */
    long long _uid; /* event ID */
    long _index; /* component index where to send this event */
    long _absorbed; /* flag set to TRUE when this event is to be removed/ignored */
    long _scattered; /* flag set to TRUE when this event has interacted with the last component instance */
    long _restore; /* set to true if neutron event must be restored */
    // user variables:
    double flag;
};

typedef struct _struct_particle _class_particle;
```

RNG state is a thread-variable contained on the _particle struct. Was earlier a global state in CPU settings

Instrument and component structs built on CPU and transferred to GPU using OpenACC pragmas at the end of

```
#ifdef USE_PGI
# include <openacc.h>
acc_attach( void*&_arm_var );
acc_attach( void*&_source_var );
acc_attach( void*&_coll2_var );
acc_attach( void*&_detector_var );
#pragma acc update device(_arm_var)
#pragma acc update device(_source_var)
#pragma acc update device(_coll2_var)
#pragma acc update device(_detector_var)
acc_attach( void*&_instrument_var );
#pragma acc update device(_instrument_var)
#endif
```



Similar “host” update in FINALLY

INITIALISE

Each component will correspond to a GPU'ified function...

+ particle-loop and logic around, also running on GPU.

Init and finalisation codes run purely CPU.

```
#pragma acc routine seq
_class_Source_simple *_class_Source_simple_trace(_class_Source_simple *_comp
, _class_particle *_particle) {
ABSORBED=SCATTERED=RESTORE=0;

#define radius (_comp->_parameters.radius)
#define yheight (_comp->_parameters.yheight)
#define xwidth (_comp->_parameters.xwidth)
#define dist (_comp->_parameters.dist)
#define focus_xw (_comp->_parameters.focus_xw)
#define focus_yh (_comp->_parameters.focus_yh)
#define E0 (_comp->_parameters.E0)
#define dE (_comp->_parameters.dE)
#define lambda0 (_comp->_parameters.lambda0)
#define dlambd a (_comp->_parameters.dlambd a)
#define flux (_comp->_parameters.flux)
#define gauss (_comp->_parameters.gauss)
#define target_index (_comp->_parameters.target_index)
#define pmul (_comp->_parameters.pmul)
#define square (_comp->_parameters.square)
#define srcArea (_comp->_parameters.srcArea)
#define tx (_comp->_parameters.tx)
#define ty (_comp->_parameters.ty)
#define tz (_comp->_parameters.tz)
SIG_MESSAGE("[_source_trace] component source=Source_simple() TRACE [Source_simple.comp:127]");
double chi,E,lambda,v,r, xf, yf, rf, dx, dy, pdir;
t=0;
z=0;

if (square == 1) {
    x = xwidth * (rand01() - 0.5);
    y = yheight * (rand01() - 0.5);
} else {
    chi=2*PI*rand01();
    r=sqrt(rand01())*radius;
    x=r*cos(chi);
    y=r*sin(chi);
}
randvec_target_rect_real(&xf, &yf, &rf, &pdir,
/* Choose point on source */
/* with uniform distribution. */
tx, ty, tz, focus_xw, focus_yh, ROT_A_CURRENT_COMP, x, y, z, 2);
.... etc
```

“Scatter-gather” approach not far from what we do in MPI settings, i.e. :

GPU case:

N particles are calculated in parallel in N GPU threads. (Leave to OpenACC/device how many actually are running at one time)

CPU case:

N particles are calculated in M serial chunks over M processors.

Contains component trace section

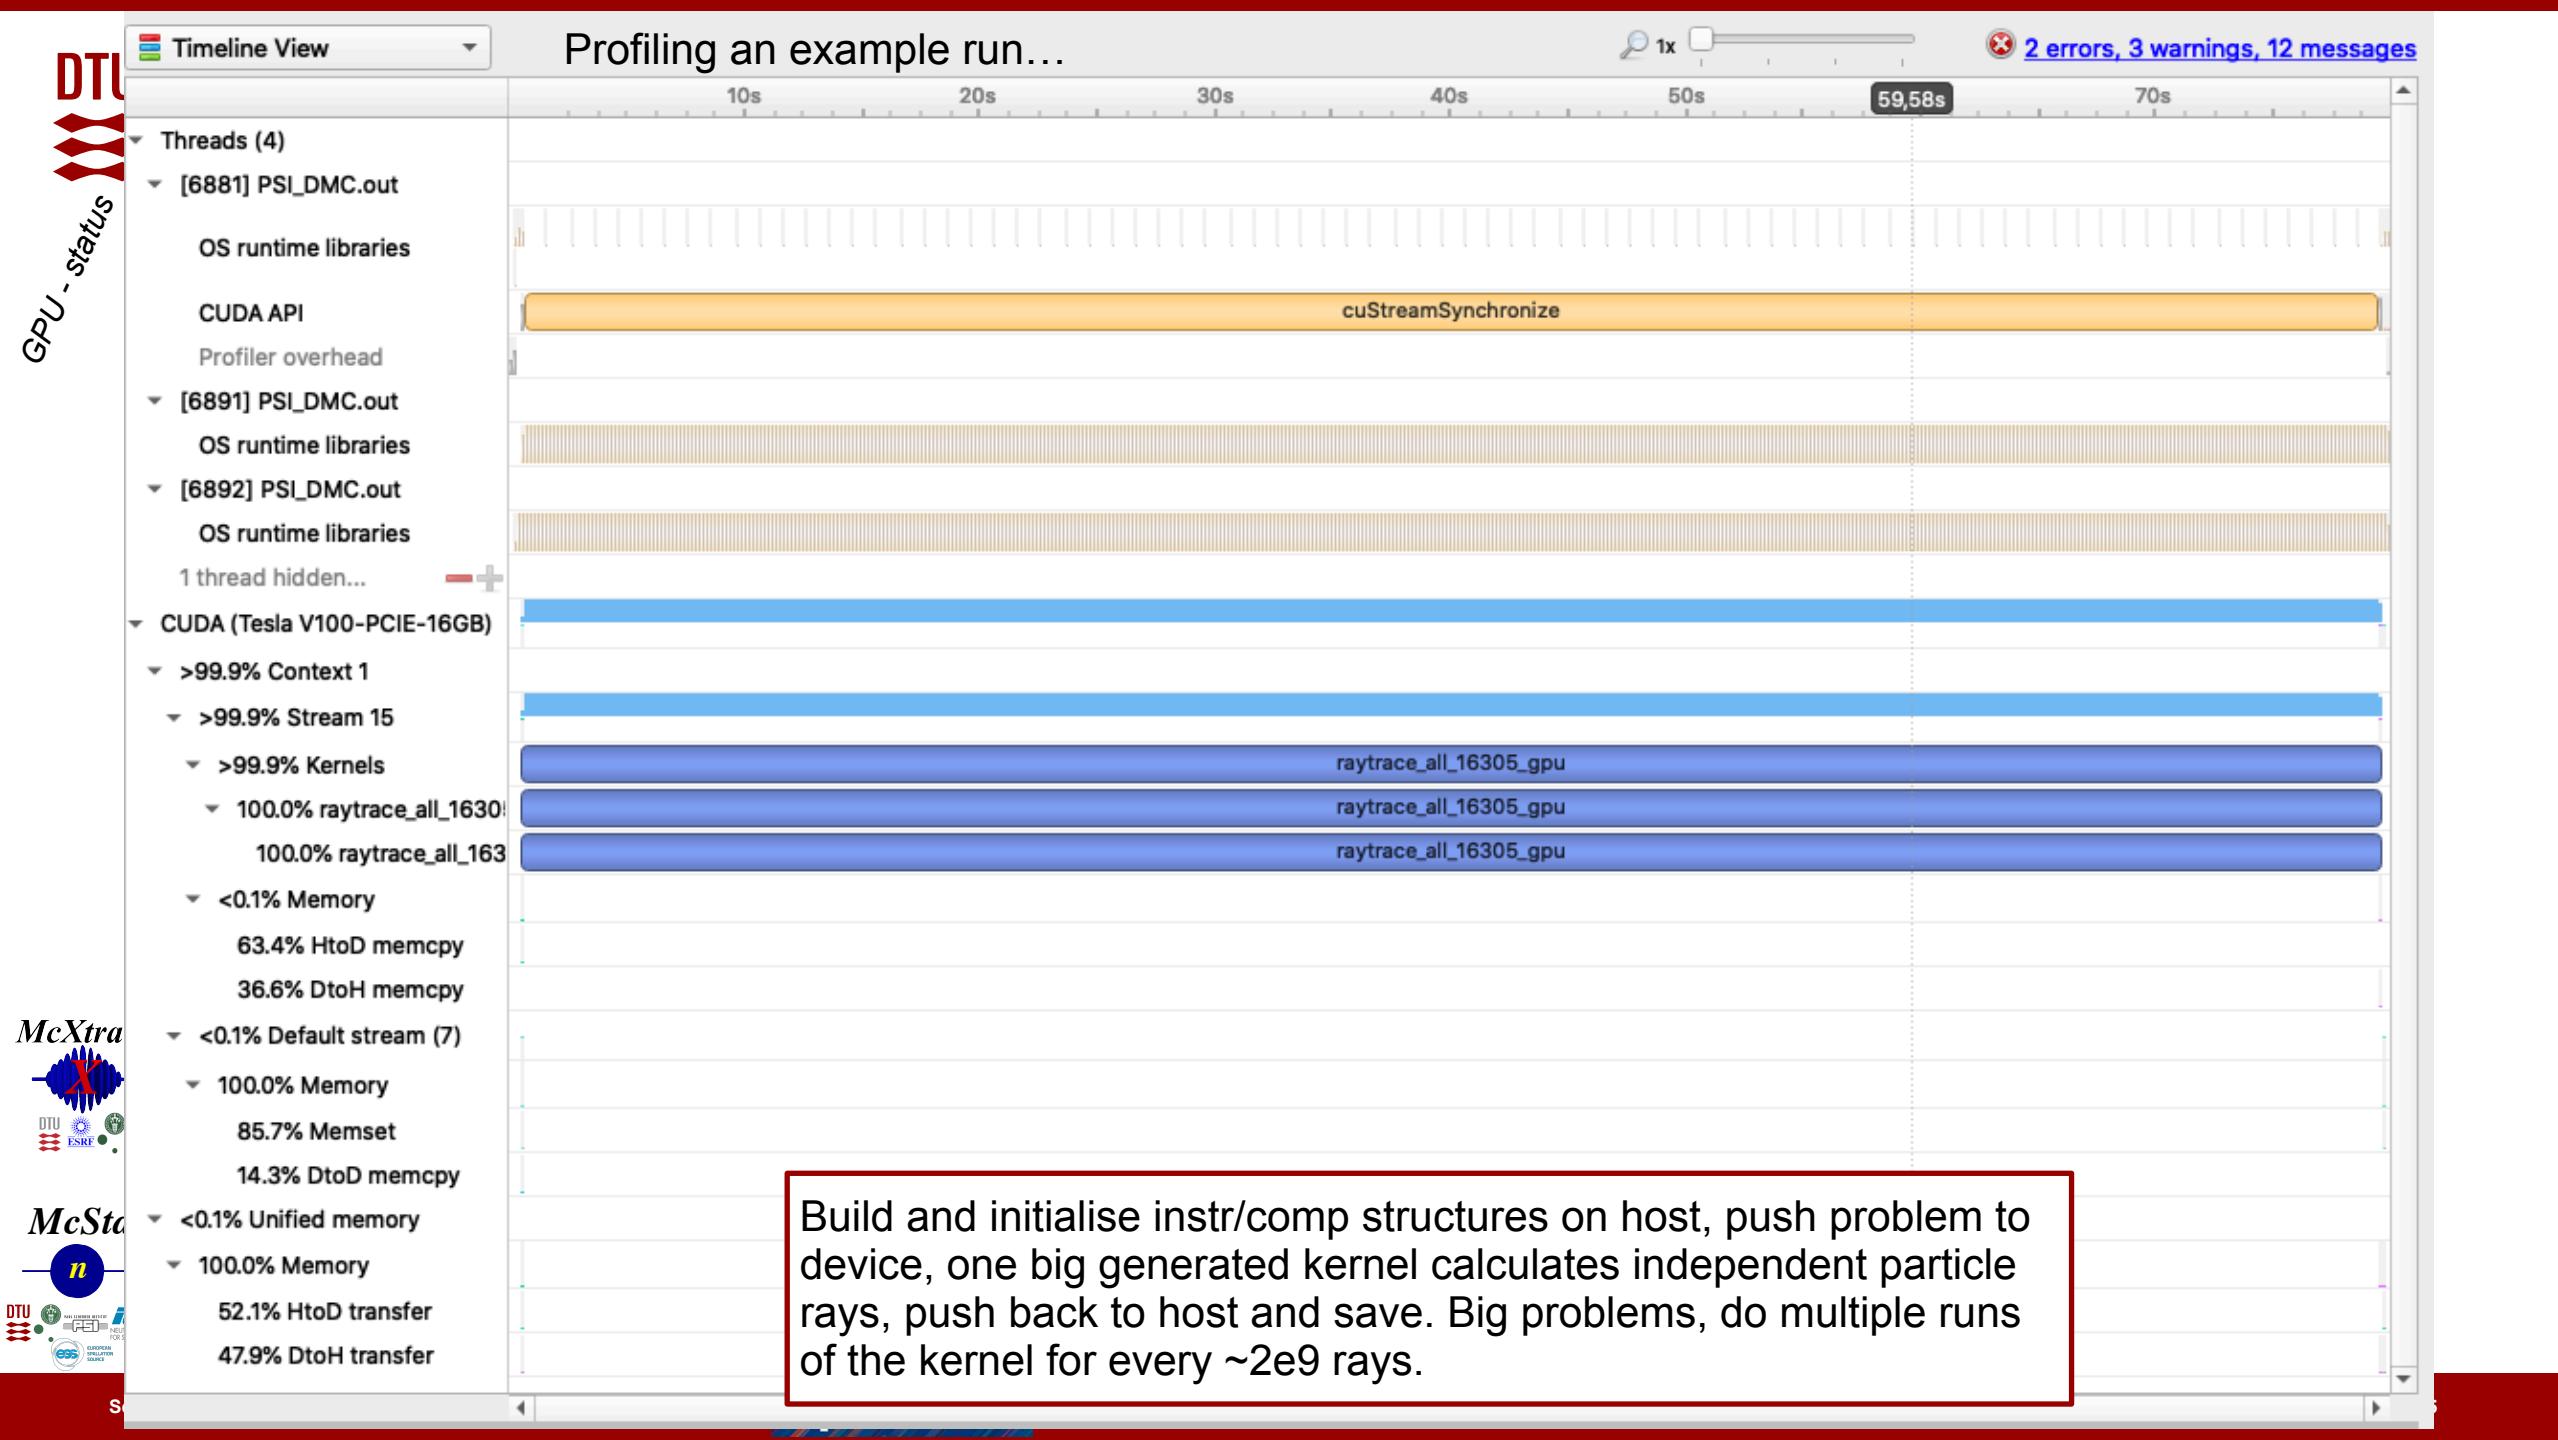
Anatomy of a McStas GPU run (*)

- Init, geometry, files etc. read on CPU
 - MPI if needed
- Memory-structures
 - Built on CPU
 - Marked for transfer to GPU (`#pragma acc declare create etc.`)
 - Initialised and synced across
 - Trace-loop is a `#pragma acc parallel loop`
 - Calculation performed entirely on GPU
 - Component structs (incl. e.g. monitor-arrays) synced across
- Finally and Save runs on CPU
 - MPI merge if needed



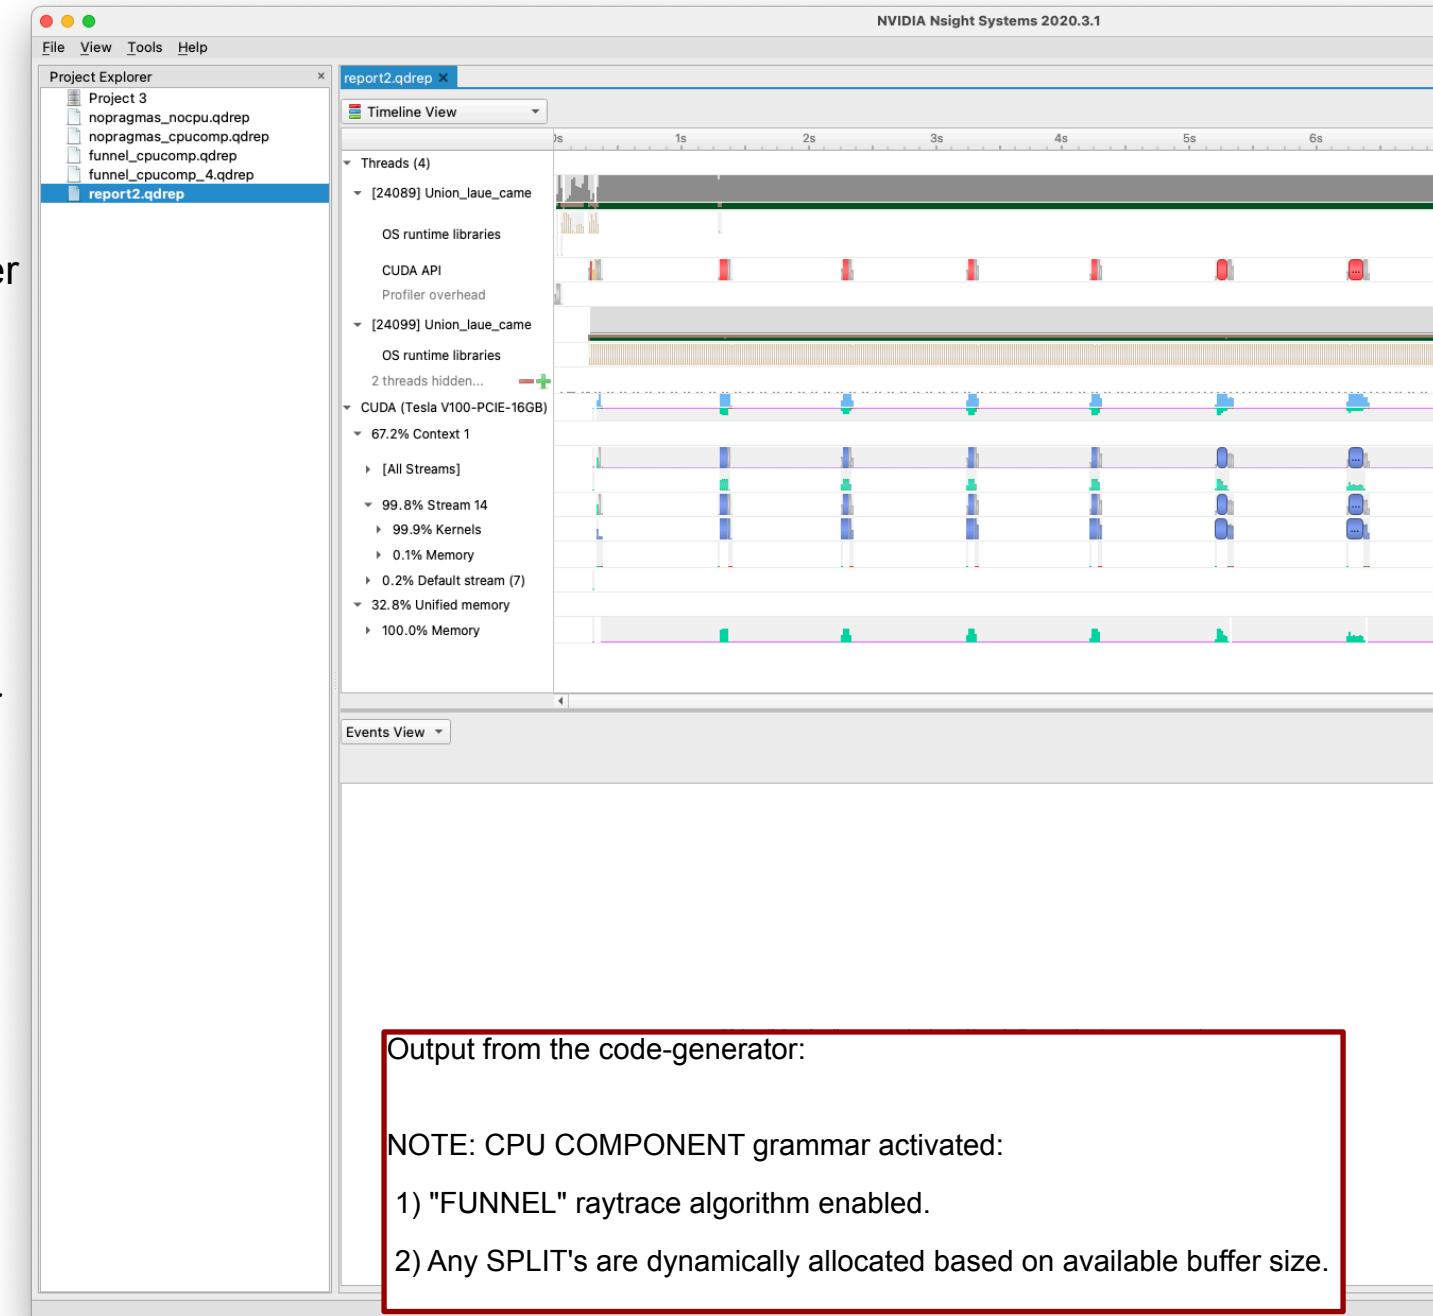
No `printf`s etc. available on GPU, automatically suppressed by `#defines`

(* Alternative layout via FUNNEL mode, see next 2 slides)



FUNNEL mode

- Activated **explicitly** using -DFUNNEL or **implicitly** using CPUCOMPONENT in instrument or NOACC in comp header
- Has N kernels / calculation zones instead of one
 - Separation at SPLIT
 - Separation if CPUCOMPONENT in instrument file
(CPUCOMPONENT A=Comp(vars=pars...))
 - Separation if a component has NOACC in the header
(See e.g. Multilayer_sample, Union_master)
- Each of these “calculation zones” is finalised before the next one initiated.
- Example:
Union: Instrument up to Union_master can be GPU, then CPU, then GPU again
 - As slow as single-core cpu... (multicore??)
 - Copying back and forth to/from GPU is costly...



Conclusion / remaining key questions

1. Why are you using OpenACC? What is your favorite feature?
 - We could **keep** most of the **core of our legacy code base intact**, some structural changes were needed. 😊 (Attempt done with OpenCL 10 y ago, no real success)
2. Did you have to use 2 diff. code bases for CPU and GPU?
 - **Nope**, just one code base, a few generated **#pragmas** and **#define's** 😊
3. Tell us about a feature you are looking for, that OpenACC lacks, why do you need it?
 - A few niche cases use function **pointers/polymorphism** on CPU, which is not available in OpenACC. Current workaround with case-hierarchy is a bit **ugly** 😐
 - **deepcopy** of 2-dimensional arrays as struct members does not seem to work correctly in our case without **managed** mode. Does not “feel” transparent wrt. possible performance penalties. 😐 We **could** fully move to 1D arrays, but this would require too much work. 😐
4. Any additional bugs, features, feedback to share?
 - We did a Hackathon together with “competing” **Python+OpenCL code RAMP** which is written from scratch. Where we tested, **we were ~40% faster**. ;-) 😊
 - In our niche setting, **curand()** turned out quite slow 😞 😐
 - **Better high-level OpenACC docs needed**. Choice of exact pragma often trial and error. 😐
 - **CUDA-mindset** / thinking seems required for **deep/full** understanding of **OpenACC**. Somewhat **contradictory** to the idea of “more science, less programming”. 😐

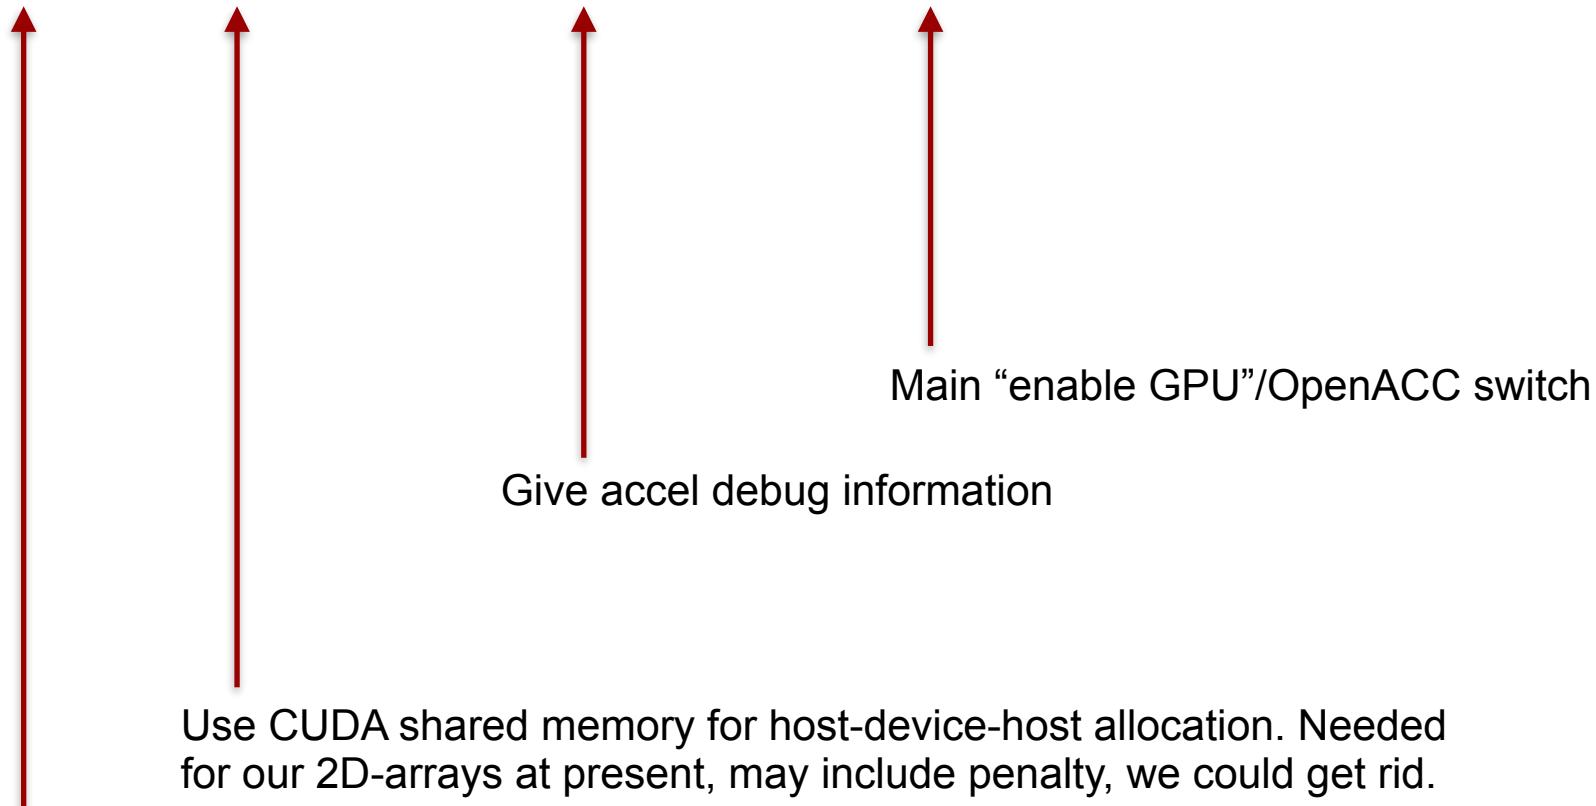

 OpenACC


FEEDBACK



Compiler settings used for GPU:

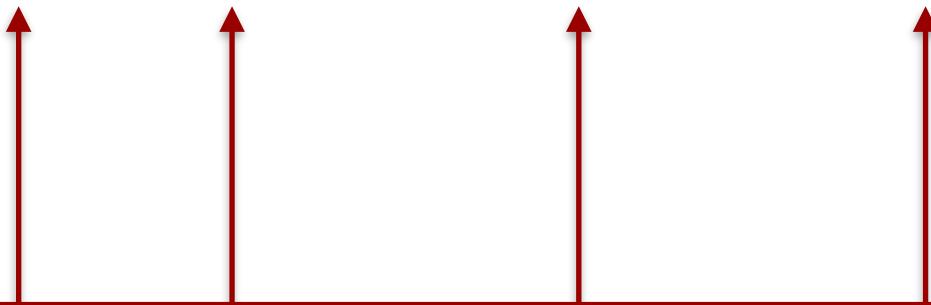
```
nvc --gpu=managed -Minfo=accel -DOPENACC
```



Generate NVIDIA GPU code, multicore also possible with ~ close to mpi performance

Compiler settings used for GPU:

```
nvc --gpu=managed -Minfo=accel -DOPENACC (-DFUNNEL)
```



Defaults for “GPU neutron loops”:
1) non-funnelled
 GPU_INNERLOOP=2147483647
2) funnelled
 GPU_INNERLOOP=1024*1024
3) control using --gpu_innerloop on binary

For CPU/threading, use below settings, with -DMULTICORE e.g. printf's are not nullified in TRACE

```
nvc --multicore -Minfo=accel -DOPENACC ( -DMULTICORE )
```

Pragmas in play...

- Data need to be transferred to the GPU, we use
 - `#pragma acc declare create(VAR) // put at VAR declaration`
`double VAR;`
 - `#pragma acc update device(VAR) // after assignment`
- Main particle loop has a
 - `#pragma acc parallel loop`
`for (unsigned long pidx=0 ; pidx < innerloop ; pidx++) {`
- Any function to be evaluated on GPU needs. Put in place by code-generator whenever we can...
 - `#pragma acc routine // on fct. prototype or on actual function def.`
`_class_Source_div *class_Source_div_trace(_class_Source_div *_comp , _class_particle *_particle) {`
- If writing to VAR is necessary, this can be done *atomically* using e.g.
 - `#pragma acc atomic`
`VAR = VAR + 1; // ++ operators etc. is too complex`
- We pull data back using this mechanism
 - `#pragma acc update host(VAR)`

We also use

- * `openacc.h` e.g. for “attaching device pointers”
- * `acclmath.h` for a `math.h` GPU replacement

+ some self-made replacements for e.g. string handling that are (otherwise) not available for GPU.

What doesn't work

- Function pointers are not available on GPU
 - Solutions:
 - Code around if possible (integration routine pr. specific function to be integrated...)
 - Mark the component NOACC (means: “trigger FUNNEL mode and run comp on CPU” with acc sections around, see later slide)
- Variadic functions are not available on GPU
- User-defined fieldfunctions for polarisation had to be abandoned
 - No solution yet, may become handled via grammar
- External libs generally can not be used in TRACE (“#pragma....” hard to add on 3rd party codes)
 - Handle in INIT / FINALLY (MCPL)
 - NOACC (GSL etc.)
- Union master is for now NOACC, will eventually become supported on GPU



Conclusions

- **It really does work nicely!**
- **Code changes** much **less invasive** than envisioned!
- It often gives a speedup of **1-2 orders** of magnitude over 1 cpu
- **Most things work**
(we have workarounds or solutions in the pipe almost all of the rest)
- **Documentation** on McCode Wiki...
- McStas 3.0 is as of yet “ported” to GPU but **not fully “optimised” performance-wise**, we will try to **go to another Hackathon < This one!**
- **Union** needs a dedicated **Hackathon**



Ideas to investigate / work on?

- Investigate **getting rid of ‘managed mode’**
 - (2D array-vars > 1D, add #pragma acc shape and make simpler structs?)
- Get “**multicore**” functional on “**NOACC**” FUNNEL sections?
- **64bit vs. 32bit Floats** (often only ~ factor 2 speed up)
 - MCNUM where applicable, use a define to decide
- **Optimise a set of selected instrument models**, for instance these:
 - Source + monitor ideal case
 - simple optics model from McStas (guide system?)
 - instrument with nested SPLITs (templateTAS?)
 - Powder diffractometer from McXtrace
 - OFF-oriented instrument, e.g. Airport_Scanner McXtrace
 - Laue instrument using Single crystal - NMX?
- Wanted outcomes:
 - **Learn to practically use the Nvidia profilers** (ncu / nvprof)
 - Get a better **understanding of the limitations** of the current implementation
 - Try to **identify obvious bottlenecks**
- Other
 - share/ snippets > **proper shared libs**? CPU conda libs, tooling for GPU libs maybe?

The team, Nvidia mentors and Hackathon hosts :-)



Vishal Metha



Christian Hundt



Alexey Romanenko



Guido Juckeland



Sebastian von Alfthan